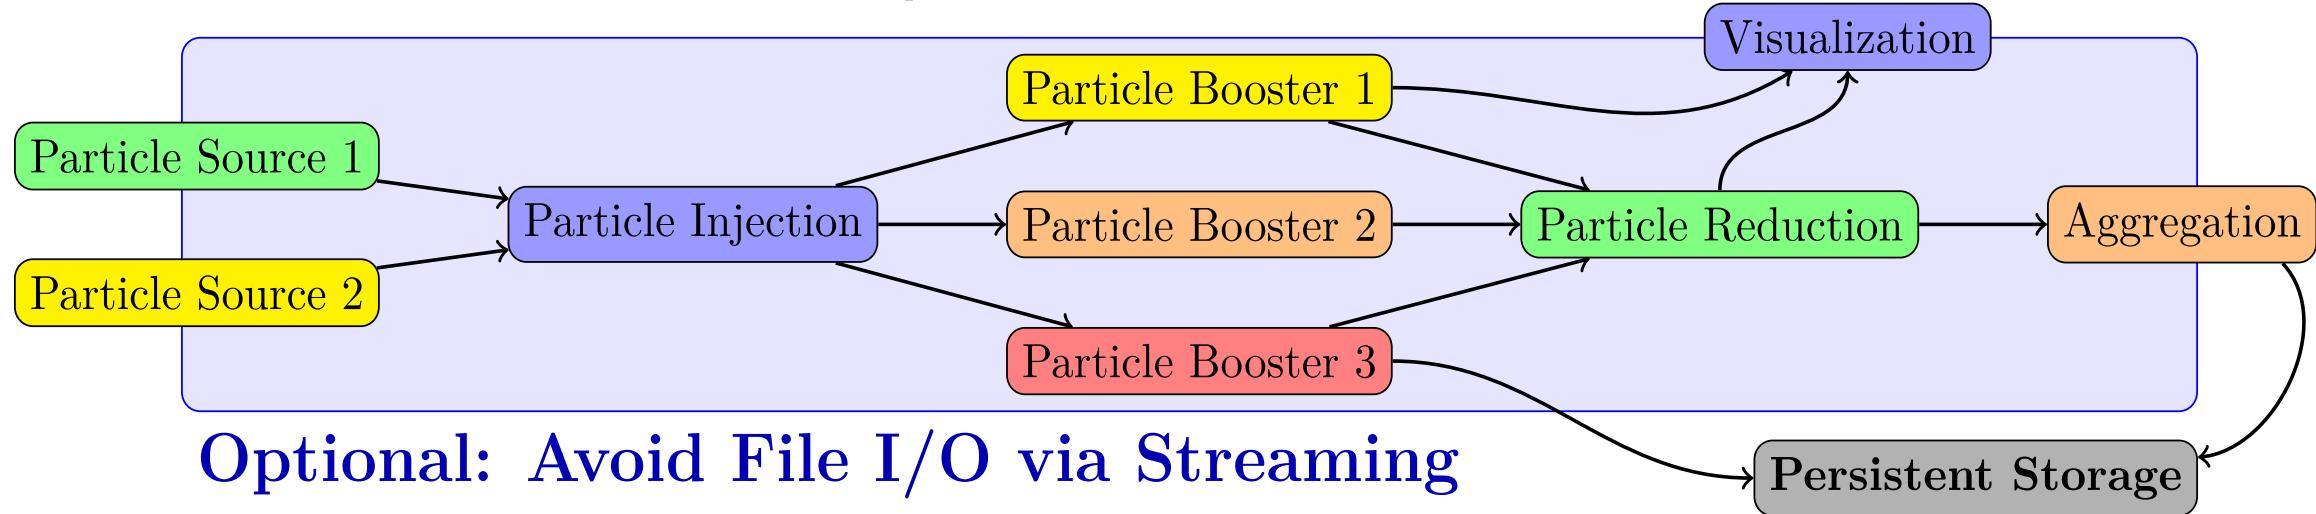


Plasma-PEPSC Workshop
23 October 2024

The Open Standard for Particle-Mesh Data

Heterogeneity through Standardized Data

Scientific workflows are complex:

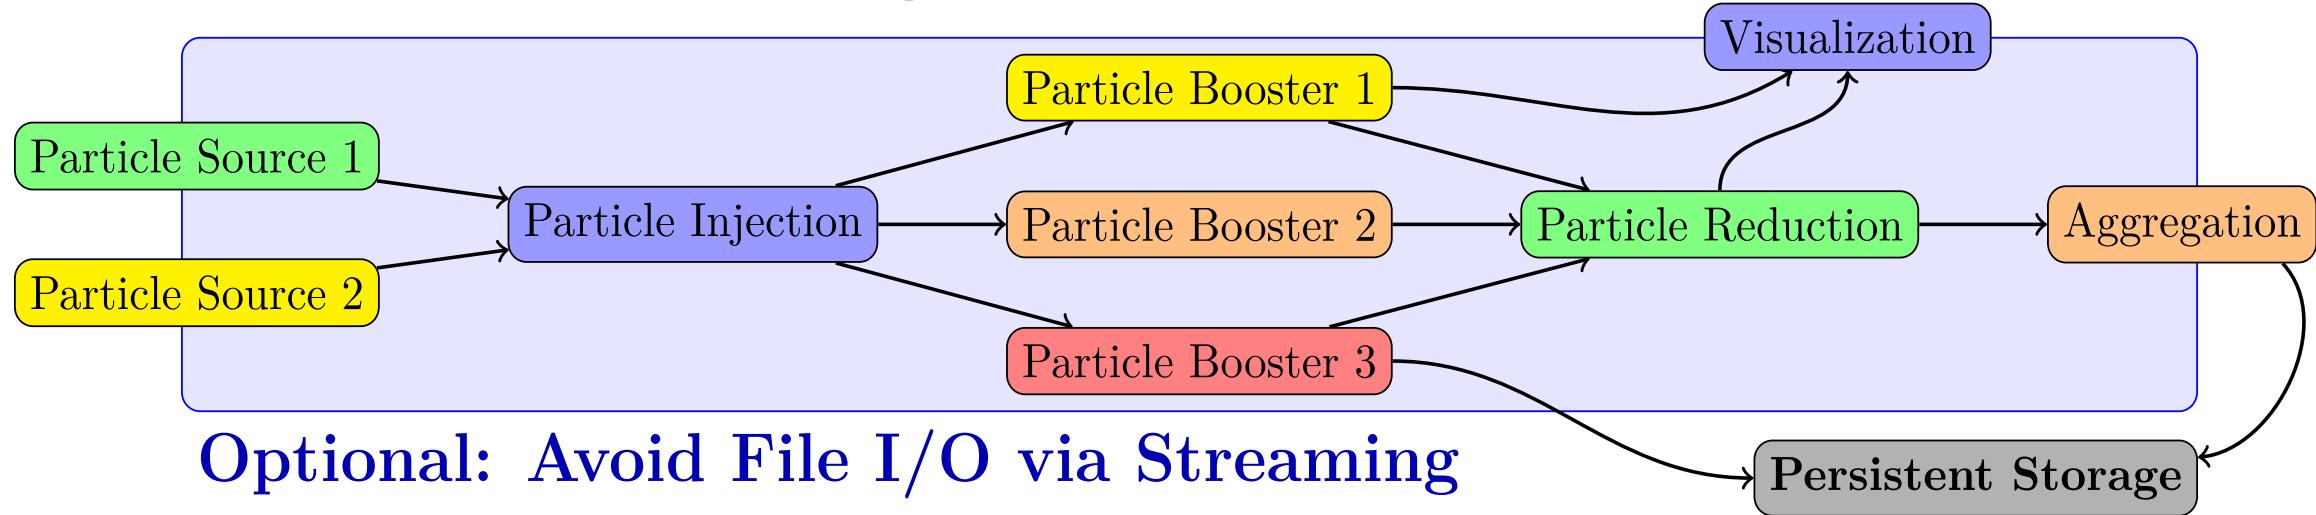


Optional: Avoid File I/O via Streaming

- need to span different **time** and **length scales**
- scientific modeling requires **multiple codes**,
collaborating in a **data processing pipeline**
- **bridge heterogeneous models** by standardization of data

Heterogeneity through Standardized Data

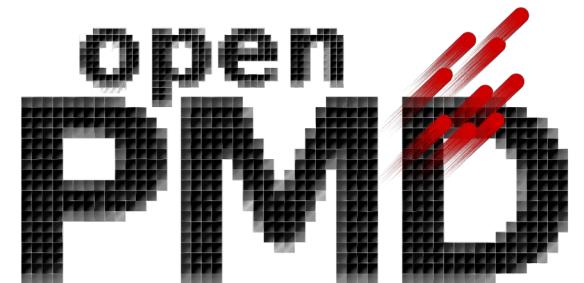
Scientific workflows are complex:



Optional: Avoid File I/O via Streaming

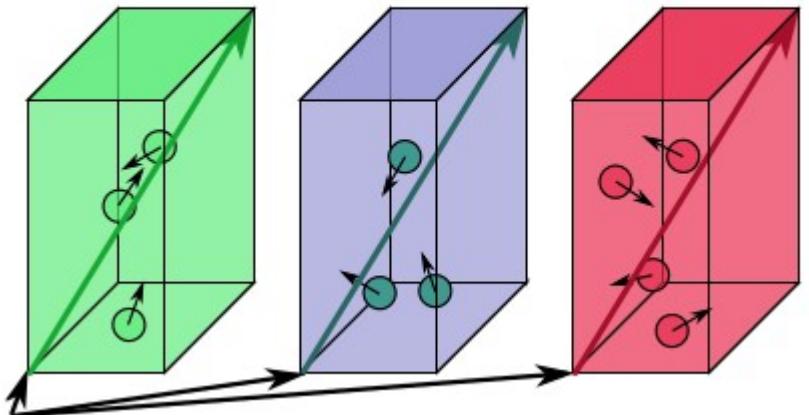
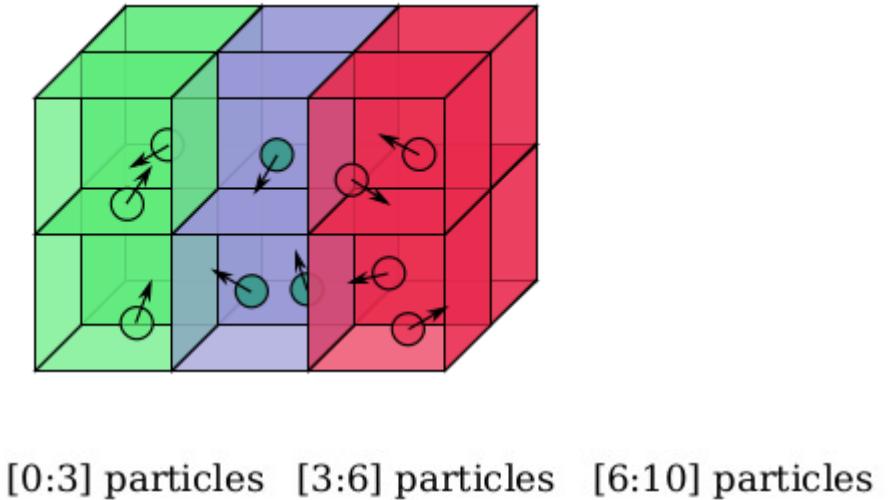


openPMD standard
for **particle-mesh data**
as communication layer



Axel Huebl et al. "openPMD: A meta data standard for particle and mesh based data". 2015. doi: 10.5281/zenodo.591699. url: <https://openPMD.org>
Franz Poeschel et al. "Transitioning from file-based HPC workflows to streaming data pipelines with openPMD and ADIOS2". 2021. doi:10.1007/978-3-030-96498-6_6

What is particle-mesh data?



Mesh

n-dimensional space,
 divided into discrete cells

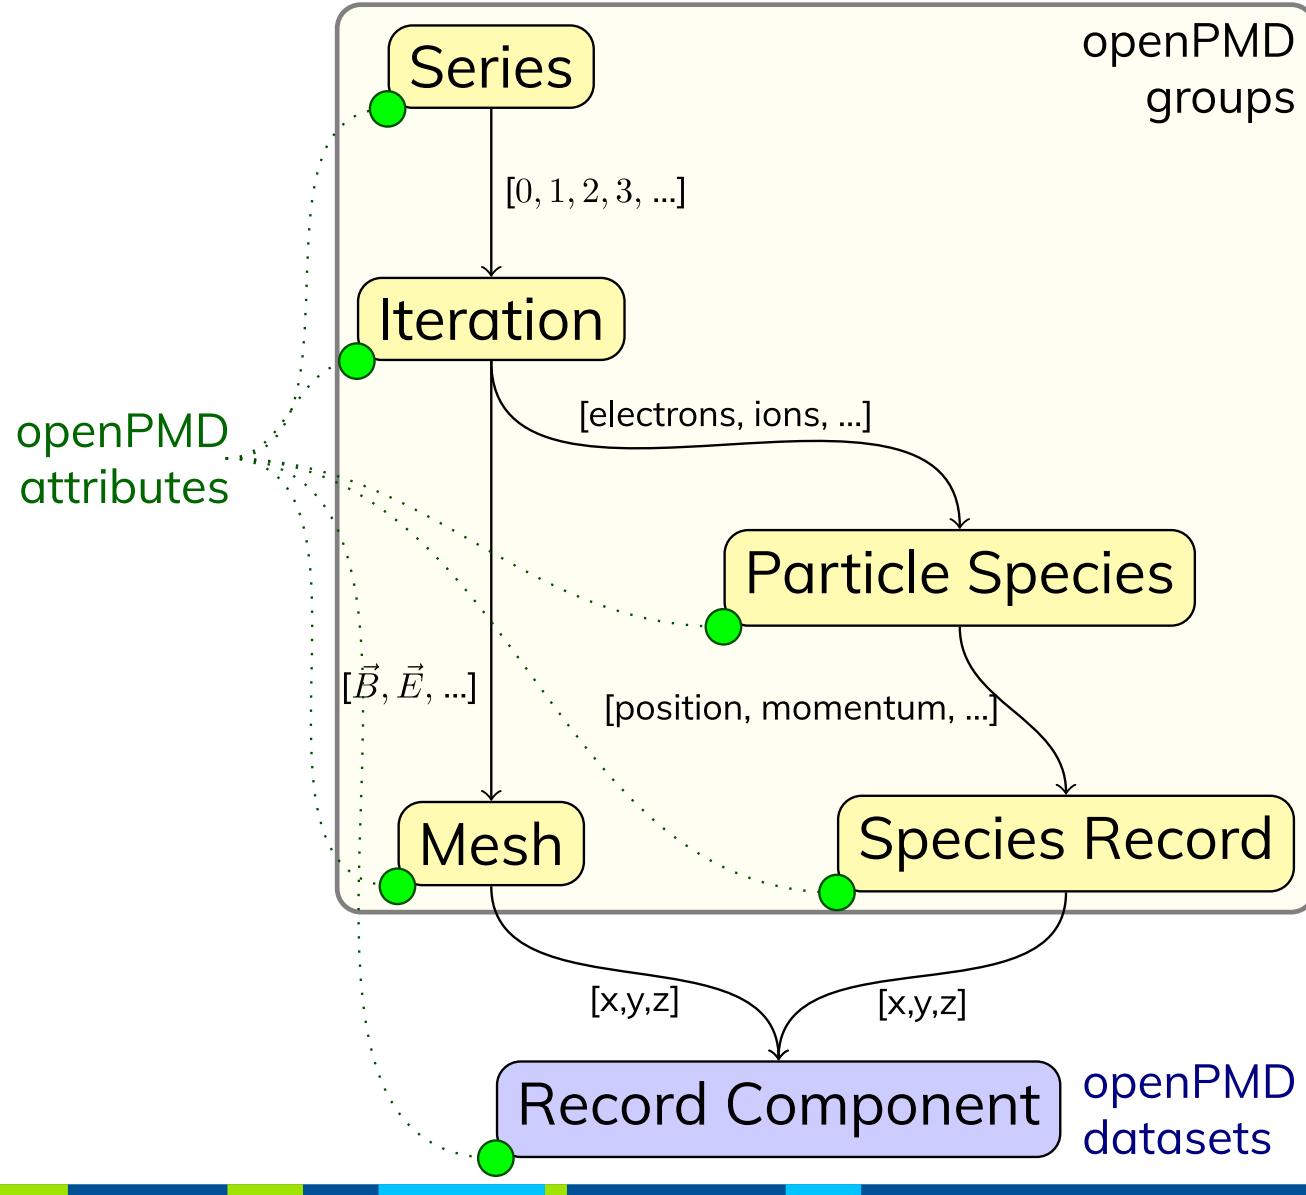
- e.g. temperature:
 store a scalar number per cell
- e.g. electrical fields:
 store a 3D vector per cell

Particles

A list of discrete objects,
 located on the mesh

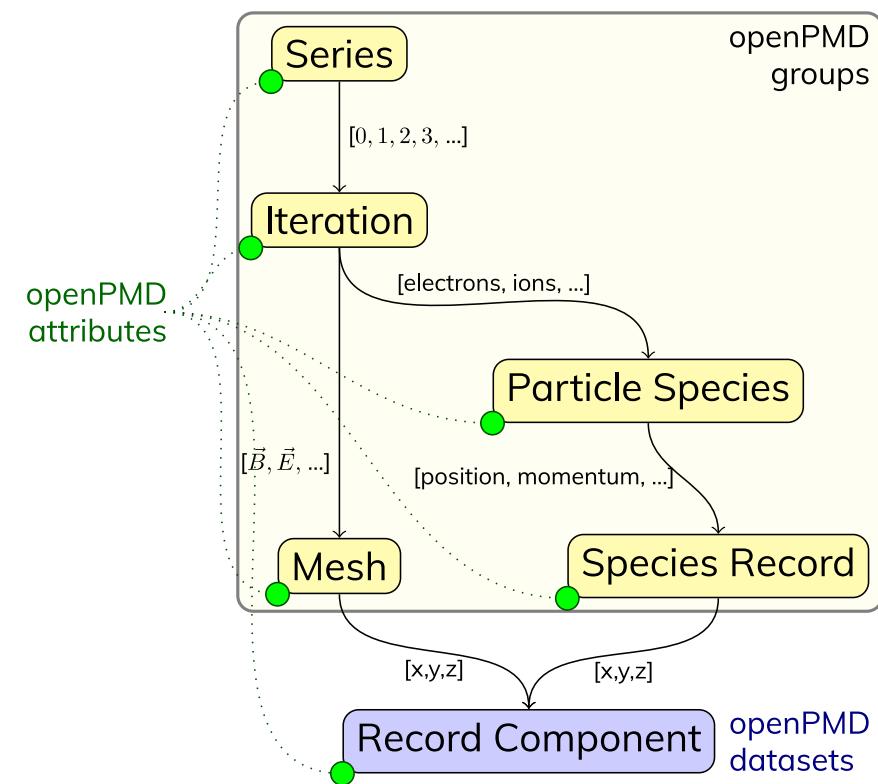
- for each particle: list its position
- optionally: list charge, weight, ...

openPMD hierarchy

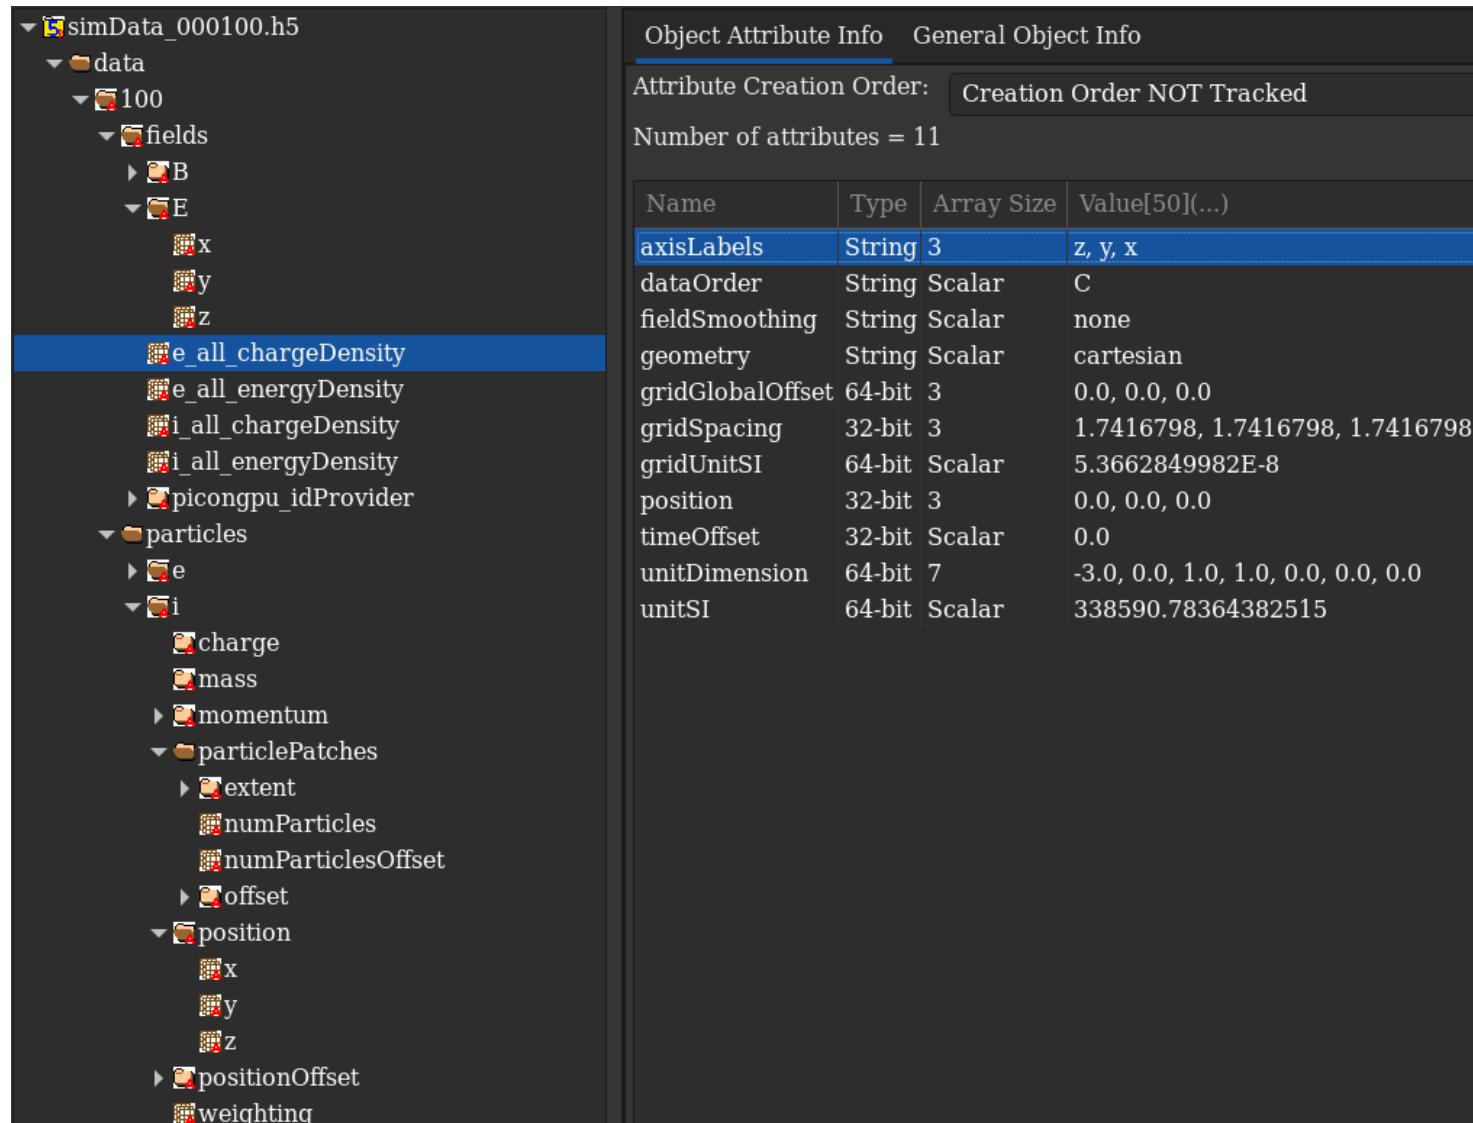


- **Structure** for series & snapshots encoded as either:
 - **files** (one file per iteration)
 - **groups** (reuse files)
 - **variables** (reuse files & variables in ADIOS2)
- Records for **physical observables** constants, mixed precision, complex numbers
- **Attributes**: unit conversion, description, relations, mesh geometry, authors, env. info, ...

Example dataset: HDF5 backend



Sample data
created with PICongPU



Example dataset: ADIOS2 backend

float	/data/50/fields/E/x	Hierarchical data organization	{128, 128, 128}	n-dim. datasets for heavyweight data
float	/data/50/fields/E/y		{128, 128, 128}	
float	/data/50/fields/E/z		{128, 128, 128}	
float	/data/50/particles/e/position/x		{50053105}	
float	/data/50/particles/e/position/y		{50053105}	
float	/data/50/particles/e/position/z		{50053105}	
int32_t	/data/50/particles/e/positionOffset/x		{50053105}	
int32_t	/data/50/particles/e/positionOffset/y		{50053105}	
int32_t	/data/50/particles/e/positionOffset/z		{50053105}	
string	/data/50/fields/E/axisLabels		attr = {"z", "y", "x"}	Attributes for self-description
string	/data/50/fields/E/dataOrder		attr = "C"	
string	/data/50/fields/E/fieldSmoothing		attr = "none"	
string	/data/50/fields/E/geometry		attr = "cartesian"	
double	/data/50/fields/E/gridGlobalOffset		attr = {0, 0, 0}	
float	/data/50/fields/E/gridSpacing		attr = {1.74168, 1.74168, 1.74168}	
double	/data/50/fields/E/gridUnitSI		attr = 5.36628e-08	
float	/data/50/fields/E/timeOffset		attr = 0	
double	/data/50/fields/E/unitDimension		attr = {1, 1, -3, -1, 0, 0, 0}	
float	/data/50/fields/E/x/position		attr = {0.5, 0, 0}	
double	/data/50/fields/E/x/unitSI		attr = 9.5224e+12	
float	/data/50/fields/E/y/position		attr = {0, 0.5, 0}	
double	/data/50/fields/E/y/unitSI		attr = 9.5224e+12	
float	/data/50/fields/E/z/position		attr = {0, 0, 0.5}	
double	/data/50/fields/E/z/unitSI		attr = 9.5224e+12	

Example dataset: JSON/TOML backend

```
{
  "attributes": {
    "basePath": "/data/%T/",
    "date": "2023-10-11 09:57:53 +0200",
    "iterationEncoding": "fileBased",
    "iterationFormat": "simData_%06T",
    "meshesPath": "fields/",
    "openPMD": "1.1.0",
    "openPMDextension": 0,
    "particlesPath": "particles/",
    "picongpuIOVersionMajor": 2,
    "picongpuIOVersionMinor": 0,
    "software": "PICConGPU",
    "softwareVersion": "0.7.0-dev"
  },
  "data": {
    "100": {
      "attributes": {
        "cell_depth": 4.252342224121094,
        "cell_height": 1.0630855560302734,
        "cell_width": 4.252342224121094,
        "dt": 1,
        "eps0": 169.19711303710938,
        "many": "more"
      },
      .....
      "fields": {
        "e_all_energyDensity": {
          "attributes": {
            "axisLabels": [ "z", "y", "x" ],
            "geometry": "cartesian",
            "gridGlobalOffset": [ 0, 0, 0 ],
            "gridSpacing": [ 4.25, 1.06, 4.25 ],
            "gridUnitSI": 4.1671151662e-08,
            "position": [ 0, 0, 0 ],
            "timeOffset": 0,
            "unitDimension": [ -1, 1, -2, 0, 0, 0, 0, 0 ],
            "unitSI": 7869098408118.734
          },
          "datatype": "FLOAT",
          "data": [
            [
              [
                [
                  "multidimensional dataset here"
                ]
              ]
            ]
          ]
        }
      }
    }
  }
}
```

- Part of the package: No need to install 3rd-party dependencies
- Useful for debugging and prototyping
- Limited parallel support
- Courtesy to Nils Lohmann's JSON library for C++
- With recent release:
Convert output to TOML
Idea: openPMD formatted configuration files

Reference Implementation in C++ & Bindings: Python and Julia



Online Documentation:
openpmd-api.readthedocs.io

The screenshot shows the documentation page for the openpmd-api. It includes a sidebar with sections for Installation, Changelog, Upgrade Guide, Usage, Concepts, and various data handling methods like First Write, Open, Iteration, Attributes, Data, Record, Units, Register Chunk, and Flush Chunk. The main content area shows how to use the API after installation, providing code snippets for C++17 and Python. The C++17 snippet uses #include <openPMD/openPMD.hpp> and the Python snippet uses import openpmd_api as io.

Open-Source Development & Tests:
github.com/openPMD/openPMD-api

The screenshot shows a GitHub Actions CI status page. It displays a green checkmark icon and the text "All checks have passed" followed by "25 successful checks". Below this, it lists five successful builds with their respective environments and durations: macOS / appleclang12_py_mpi_h5_ad2 (pull_request) Successful in 17m, Windows / MSVC w/o MPI (pull_request) Successful in 6m, Intel / ICC C++ only (pull_request) Successful in 7m, Tooling / Clang ASAN UBSAN (pull_request) Successful in 58m, and Nvidia / CTK@11.2 (pull_request) Successful in 4m. Each entry has a "Details" link.

Rapid and easy installation on any platform:



```
python3 -m pip install  
openpmd-api
```



```
conda install  
-c conda-forge  
openpmd-api
```



```
brew tap openpmd/openpmd  
brew install openpmd-api
```



```
spack install  
openpmd-api
```



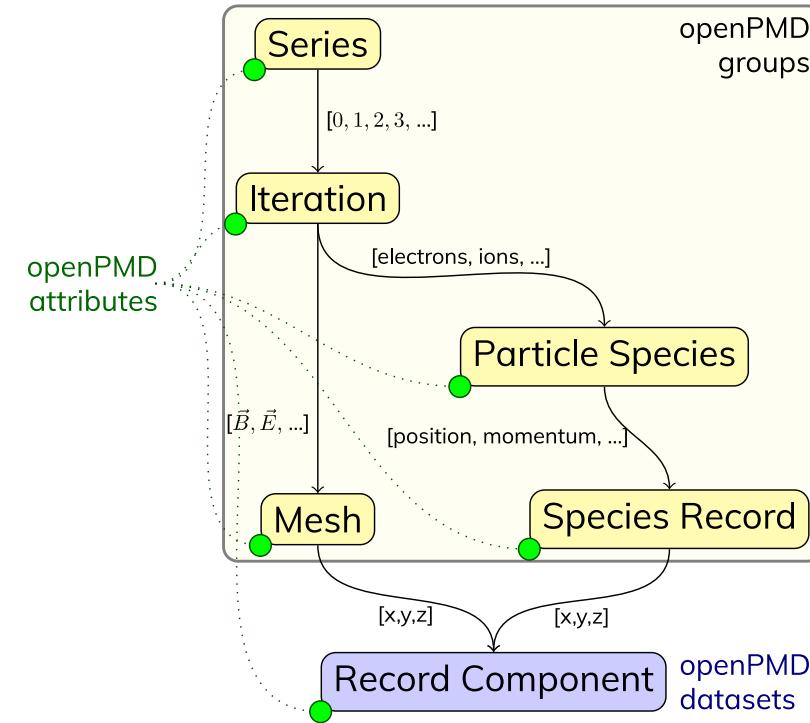
```
cmake -S . -B build  
cmake --build build  
--target install
```



```
module load openpmd-api
```

A Huebl, F Poeschel, F Koller, J Gu, et al.
"openPMD-api: C++ & Python API for Scientific I/O with openPMD" (2018) DOI:10.14278/rodare.27

Hands-On: openPMD-api: basic object model



Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Read the TODO comments inside `src/openPMDOutput.hpp`:

unitDimension

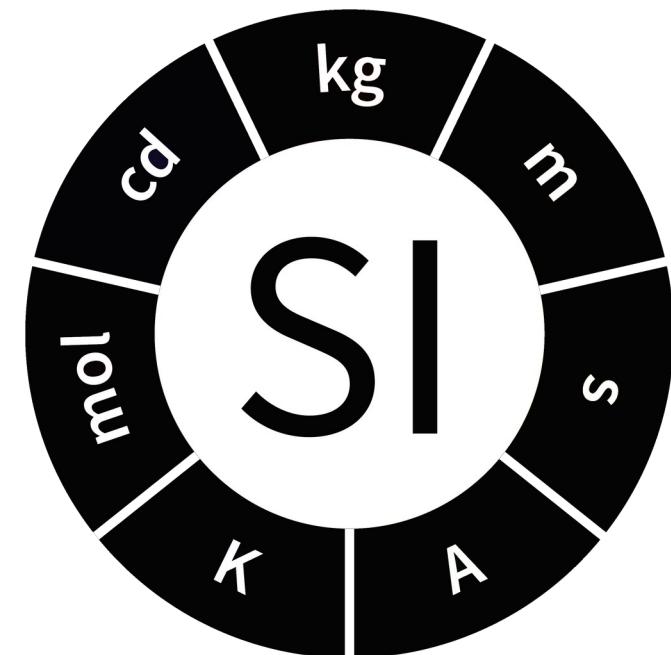
automated description of physical dimension
only powers of base dimensions

length **L**, mass **M**, time **T**, electric current **I**,
thermodynamic temperature **theta**,
amount of substance **N**, luminous intensity **J**

Magnetic field: $[B] = M / (I * T^2)$
 $\rightarrow (0, 1, -2, -1, 0, 0, 0)$

unitSI (recommended)

relation to an absolute unit system

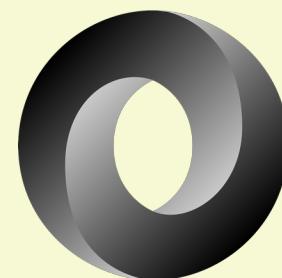


Wikimedia Commons

Findable: Standardized metadata to identify the data producer

```
string      /author           attr   = "franz"  
string      /software          attr   = "PIConGPU"  
string      /softwareVersion    attr   = "0.5.0-dev"
```

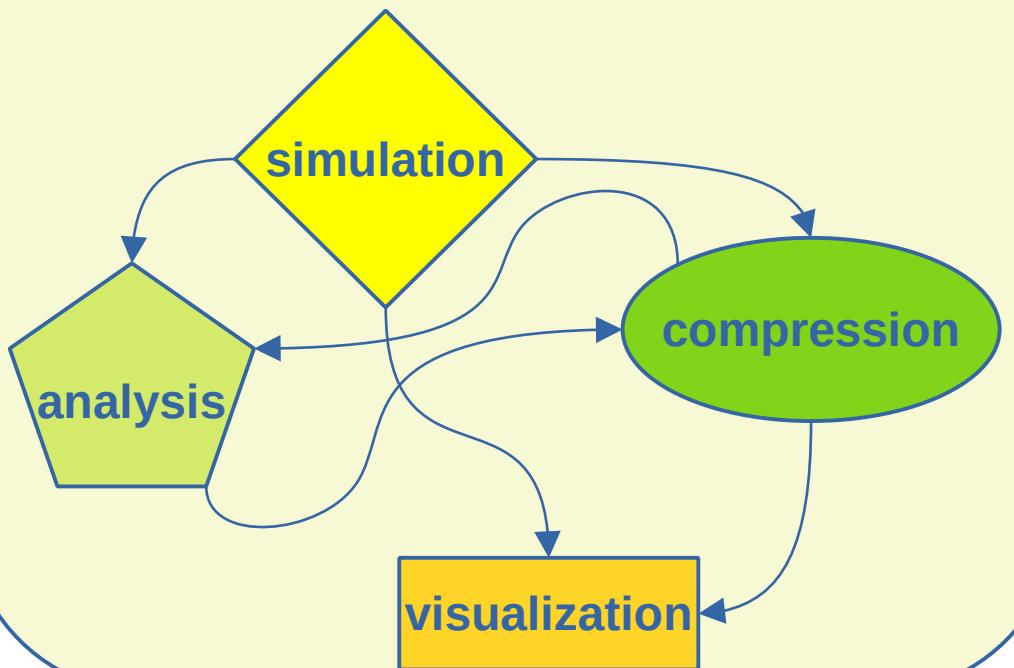
Accessible: Open standard, implementable in various formats



*currently implemented,
but not limited to

Interoperable:

Data exchange spans applications, platforms and teams



Reusable:

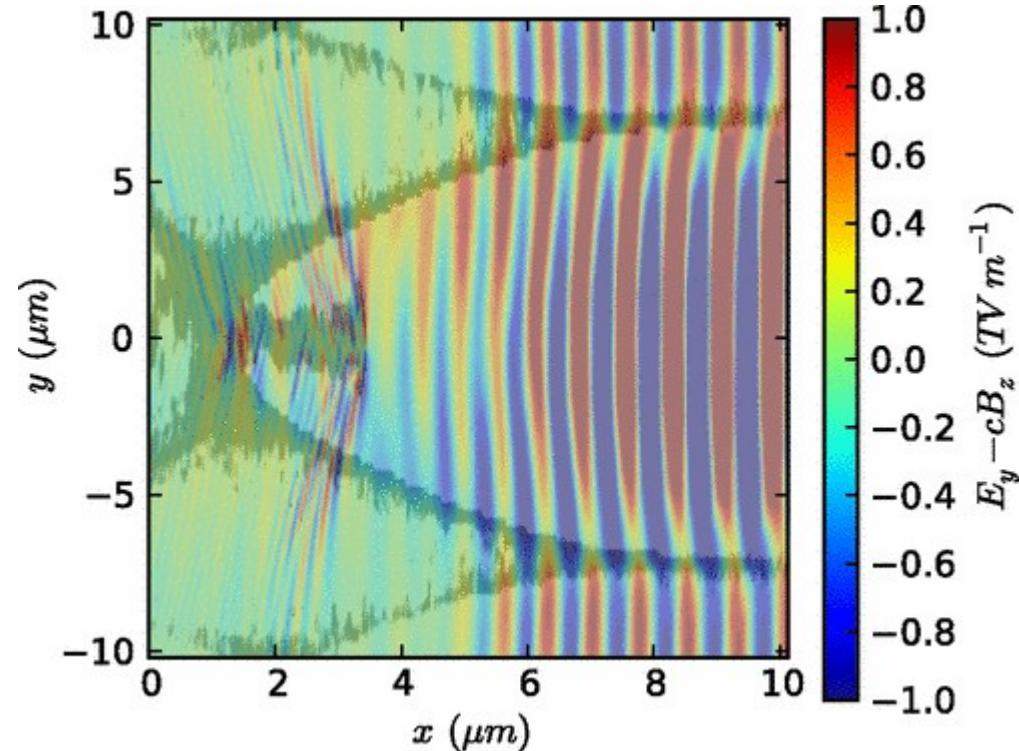
Rich and standardized description for physical quantities

Name	Value
axisLabels	[b'z' b'y' b'x']
dataOrder	b'C'
fieldSmoothing	b'none'
geometry	b'cartesian'
gridGlobalOffset	[0. 0. 0.]
gridSpacing	[4.252342 1.0630856 4.252342]
gridUnitSI	4.1671151662e-08
position	[0. 0. 0.]
timeOffset	0.0
unitDimension	[-3. 0. 1. 1. 0. 0. 0.]
unitSI	15399437.98944343

"The FAIR Guiding Principles for scientific data management and stewardship" (Mark D. Wilkinson et al.)

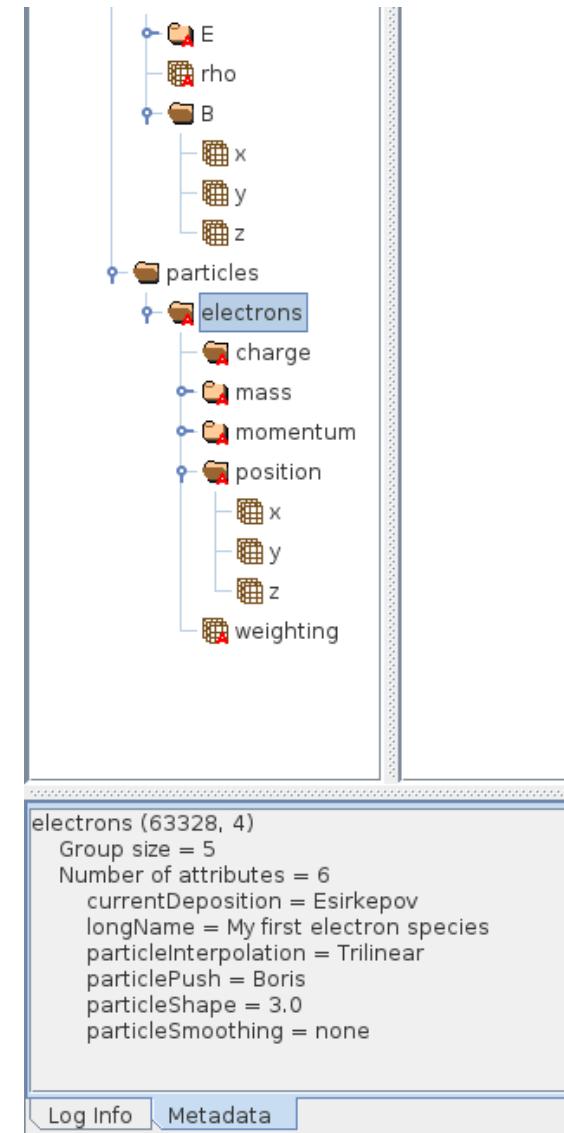
Extensions: e.g. ED-PIC

Field image → field solver, smoothing



similar:

Emittance → particle push, field solver, shape



Modeling Mesh data

```

string      /data/1000/fields/E/axisLabels           attr   = {"z", "y", "x"}
string      /data/1000/fields/E/fieldSmoothing     attr   = "none"
string      /data/1000/fields/E/geometry            attr   = "cartesian"
double     /data/1000/fields/E/gridGlobalOffset    attr   = {0, 0, 0}
float      /data/1000/fields/E/gridSpacing        attr   = {4.25234, 1.06309, 4.25234}
double     /data/1000/fields/E/gridUnitSI          attr   = 4.16712e-08
float      /data/1000/fields/E/timeOffset          attr   = 0
double     /data/1000/fields/E/unitDimension       attr   = {1, 1, -3, -1, 0, 0, 0}
float     /data/1000/fields/E/x                  {192, 1536, 192} = 0 / 0
float      /data/1000/fields/E/x/position         attr   = {0.5, 0, 0}
double     /data/1000/fields/E/x/unitSI          attr   = 1.22627e+13
float     /data/1000/fields/E/y                  {192, 1536, 192} = 0 / 0
float      /data/1000/fields/E/y/position         attr   = {0, 0.5, 0}
double     /data/1000/fields/E/y/unitSI          attr   = 1.22627e+13
float     /data/1000/fields/E/z                  {192, 1536, 192} = 0 / 0
float      /data/1000/fields/E/z/position         attr   = {0, 0, 0.5}
double     /data/1000/fields/E/z/unitSI          attr   = 1.22627e+13

float     /data/1000/fields/e_all_energyDensity   {192, 1536, 192} = 0 / 0
string      /data/1000/fields/e_all_energyDensity/axisLabels attr   = {"z", "y", "x"}
string      /data/1000/fields/e_all_energyDensity/fieldSmoothing attr   = "none"
string      /data/1000/fields/e_all_energyDensity/geometry      attr   = "cartesian"
double     /data/1000/fields/e_all_energyDensity/gridGlobalOffset attr   = {0, 0, 0}
float      /data/1000/fields/e_all_energyDensity/gridSpacing   attr   = {4.25234, 1.06309, 4.25234}
double     /data/1000/fields/e_all_energyDensity/gridUnitSI    attr   = 4.16712e-08
float      /data/1000/fields/e_all_energyDensity/position      attr   = {0, 0, 0}
float      /data/1000/fields/e_all_energyDensity/timeOffset     attr   = 0
double     /data/1000/fields/e_all_energyDensity/unitDimension  attr   = {-1, 1, -2, 0, 0, 0, 0}
double     /data/1000/fields/e_all_energyDensity/unitSI         attr   = 7.8691e+12
  
```

Mesh attributes,
including fieldSmoothing
defined by ED-PIC extension

x/y/z components for vector-type field (struct-of-array)
including component-specific metadata

Scalar-type field
x/y/z layer is skipped

Hands-On: **openPMD-api: metadata**

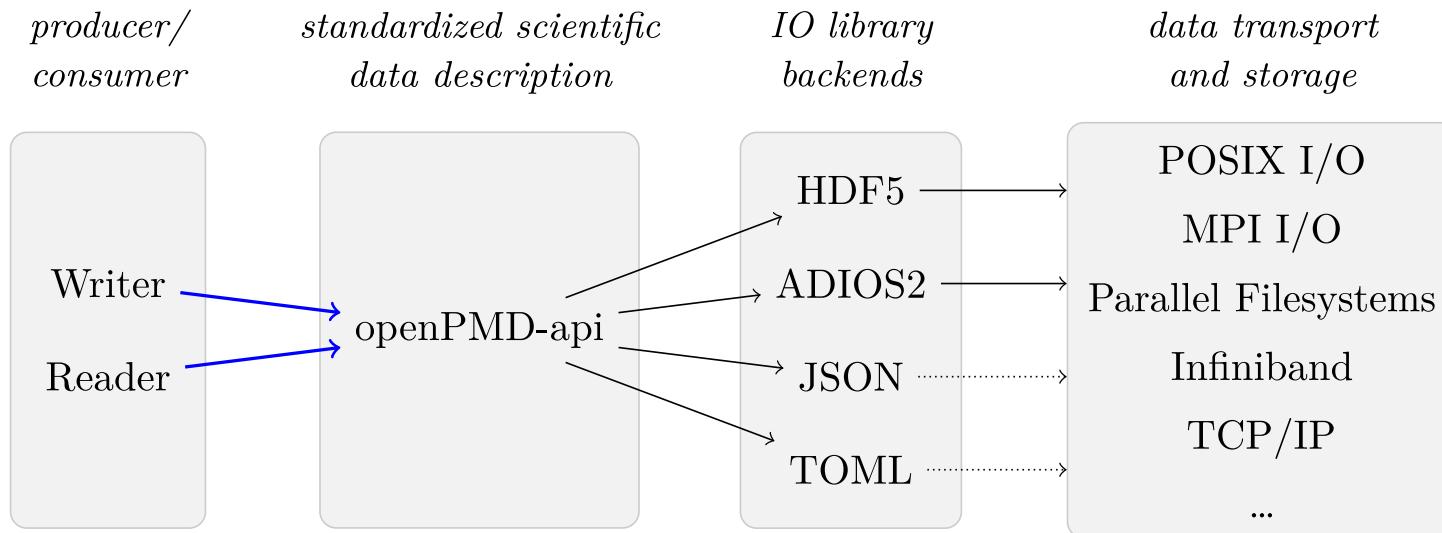
Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Read the TODO comments inside `src/openPMDOutput.hpp`:



openPMD-api – open stack for scientific I/O



```
import openpmd_api as io

# pick and configure backend via JSON/TOML or inferred from filename extension
adios_config = """
  backend = "adios2"
  [[adios2.dataset.operators]]
  type = "blosc" # activate compression
"""

mode = io.Access.create
series = io.Series("simOutput.h5", mode
                  """{"hdf5": {"vfd": {"type": "subfiling"}}}""")
series = io.Series("simOutput.bp5", mode, adios_config)
series = io.Series("simOutput.sst", mode, "@./or/load/config/from/file.json")
series = io.Series("simOutput.json", mode)
```

- MPI support at all levels
- Implemented in C++17
- Bindings in C++17, Python and (dev version only) Julia
- Specify backend at runtime: I/O library, transport, compression, streaming, aggregation, ...

Hands-On: **openPMD-api:** visualization, backend configuration

Module environment at `/project/project_465001310/workshop_software/env.sh`

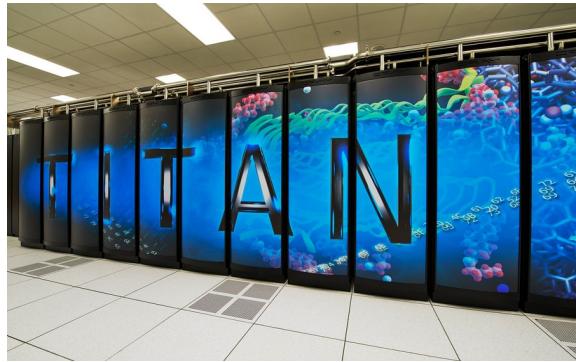
Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the instructions inside `src/next_steps.md`:



Compute Performance Outpaces Storage Performance



	Titan
Peak Performance:	27 Pflop/s
FS Throughput:	1 TiByte/s
FS Capacity:	27 PiByte

	Summit
	200 Pflop/s
	2.5 TiByte/s
	250 PiByte

	Frontier
	1.6 Eflop/s
	5~10 TiByte/s
	500~1000 PiByte

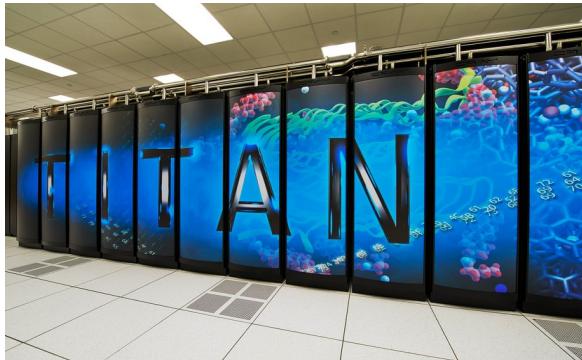
Growth Factor
~60
5~10
18~37

- **parallel bandwidth** insufficient for HPC at full scale
- **filesystem capacity** insufficient for HPC at full scale

Same trend in **experiments?**

- Increasing **camera resolutions and data rates**

Compute Performance Outpaces Storage Performance



	Titan
Peak Performance:	27 Pflop/s
FS Throughput:	1 TiByte/s
FS Capacity:	27 PiByte

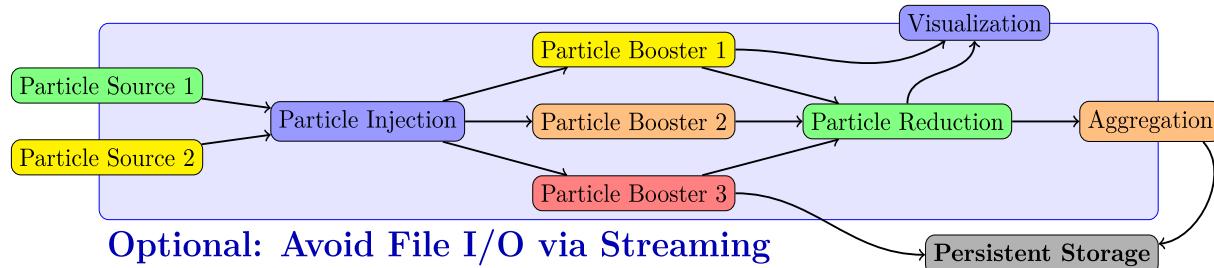
	Summit
200 Pflop/s	
2.5 TiByte/s	
250 PiByte	

	Frontier
1.6 Eflop/s	
5~10 TiByte/s	
500~1000 PiByte	

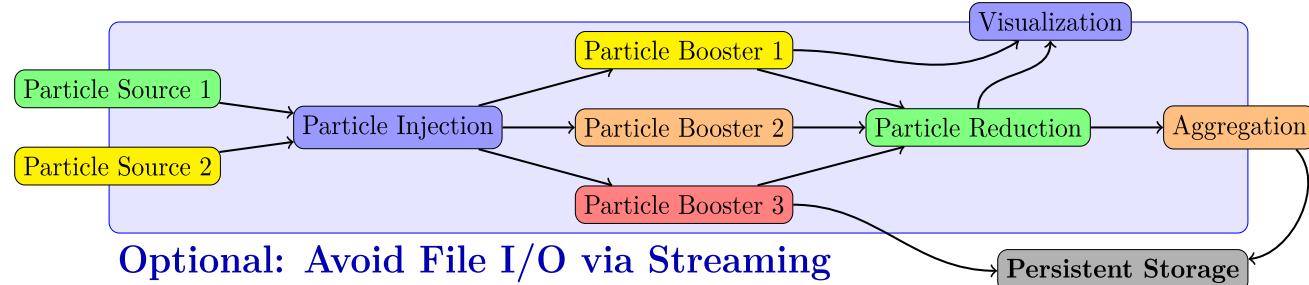
Growth Factor
~60
5~10
18~37

Why does this concern us?

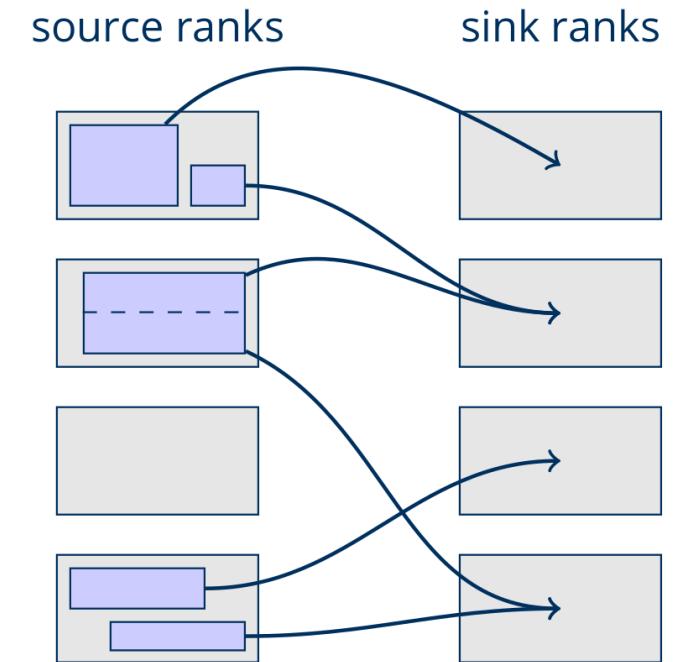
- Heterogeneous data processing pipelines traditionally have large I/O usage
- Scalable alternative: Streaming



Streaming: Don't touch the Filesystem at all

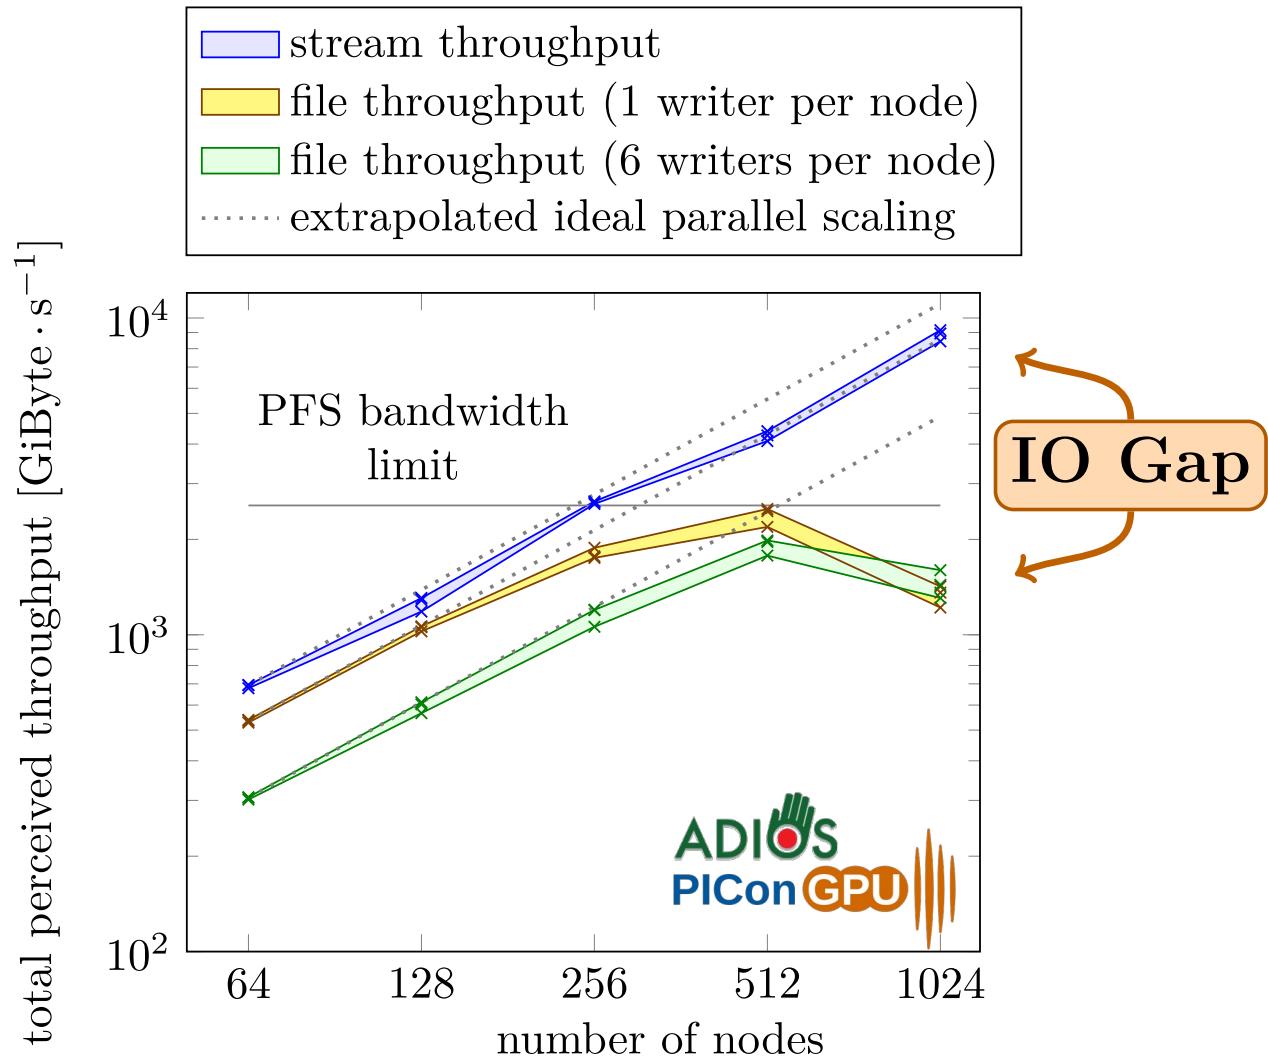


- Data processing pipelines and increasingly experiments setups have large I/O usage
- Scalable alternative: **Streaming**
e.g. via Infiniband (on HPC systems)
or wide area networks (in lab settings)



Challenge:
Compute a balanced, aligned, local mapping between two applications that remains useful in the problem domain

Break through Filesystem Bandwidth with Streaming

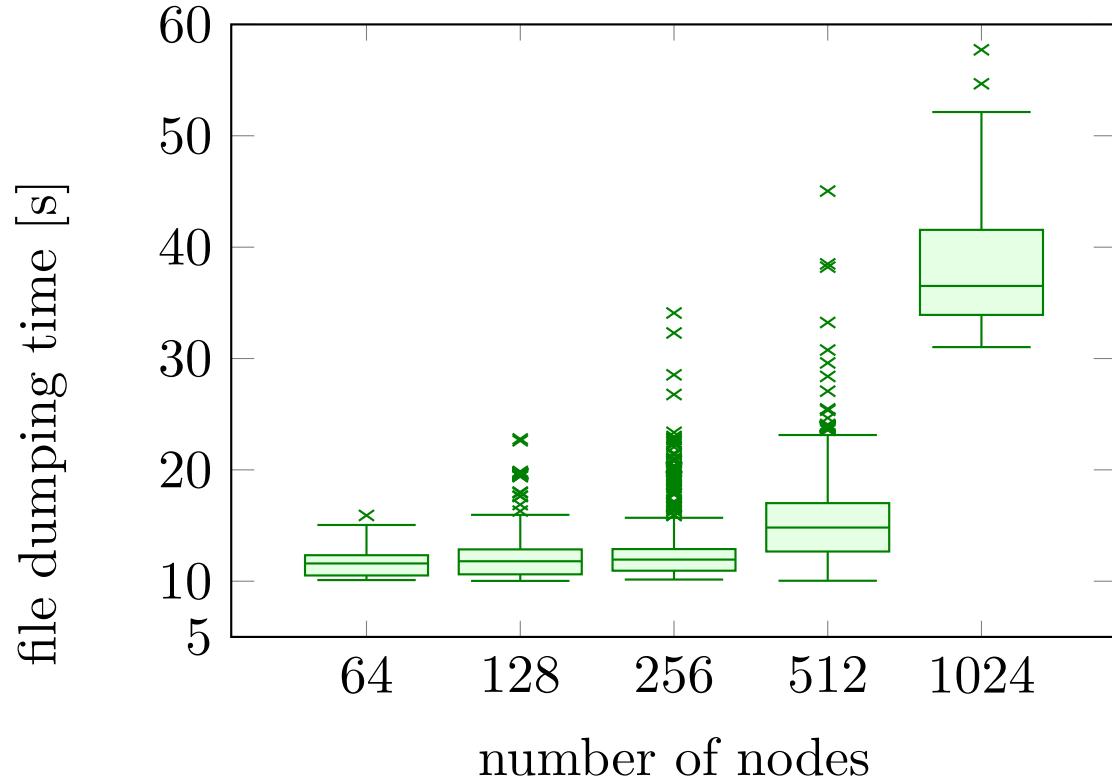


Memory-bound simulations reach the I/O system limits at a fraction of full scale

- Summit FS bandwidth (2.5TiByte/s) reached at 512 nodes (~11% of system size)
- Streaming workflows unaffected by filesystem bandwidth, use Infiniband hardware to scale beyond it

(benchmarks at 1024 nodes done after Summit system upgrade)

Summit: Performance fluctuations on single ranks

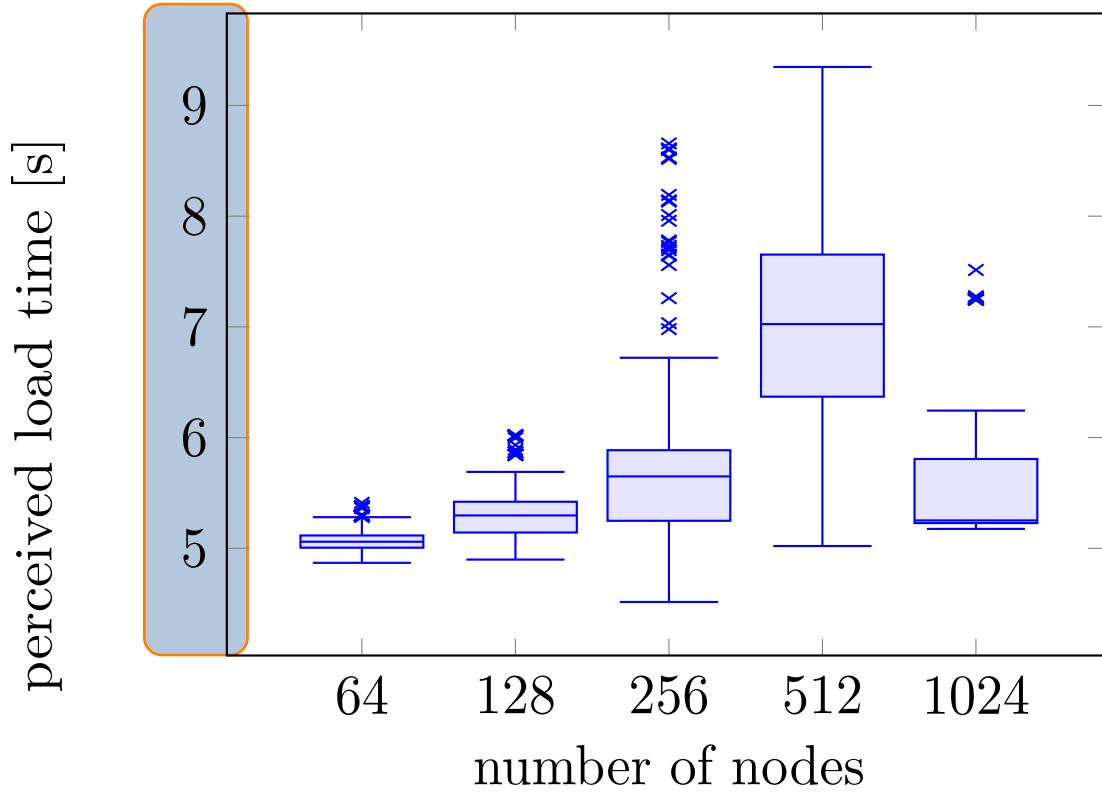


Same results, different display:

- Plot every single measurement
- Visualize reproducibility
- Box: 50% of measurement points
- Whiskers: “normal” measurements
- Others: outliers

Evaluation for file-only setup:

- Median time slightly raised at 512 nodes
Scaling stops due to PFS limit after that
- Outliers increasing with scale
- Outliers fatal in parallel contexts



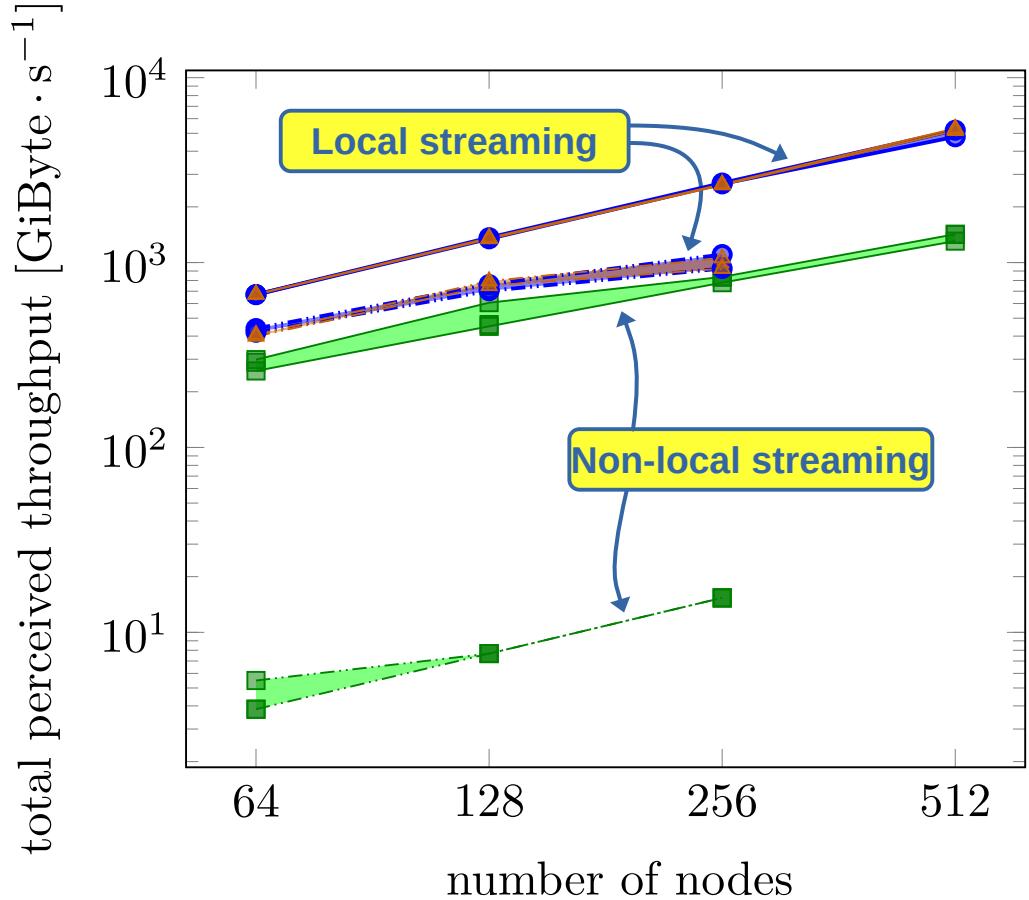
Same results, different display:

- Plot every single measurement
- Visualize reproducibility
- Box: 50% of measurement points
- Whiskers: “normal” measurements
- Others: outliers

stream+file setup (stream part):

- Overall times are lower
- Median between 5 and 7 seconds
- Outliers less dominating by far

For good throughput: Local streaming patterns, Infiniband/RDMA



Local streaming:

Distribute data chunks
only within a node

Non-local streaming:

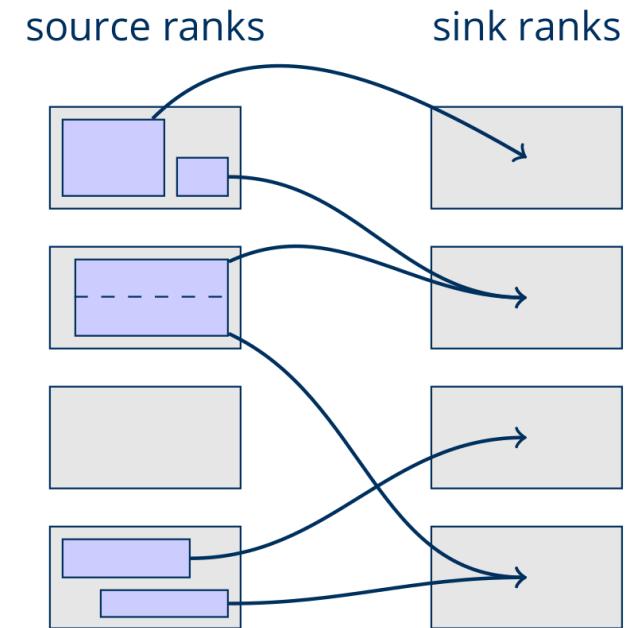
Distribute data chunks
globally, optimize for
balance and alignment

Straight lines:

Infiniband/RDMA

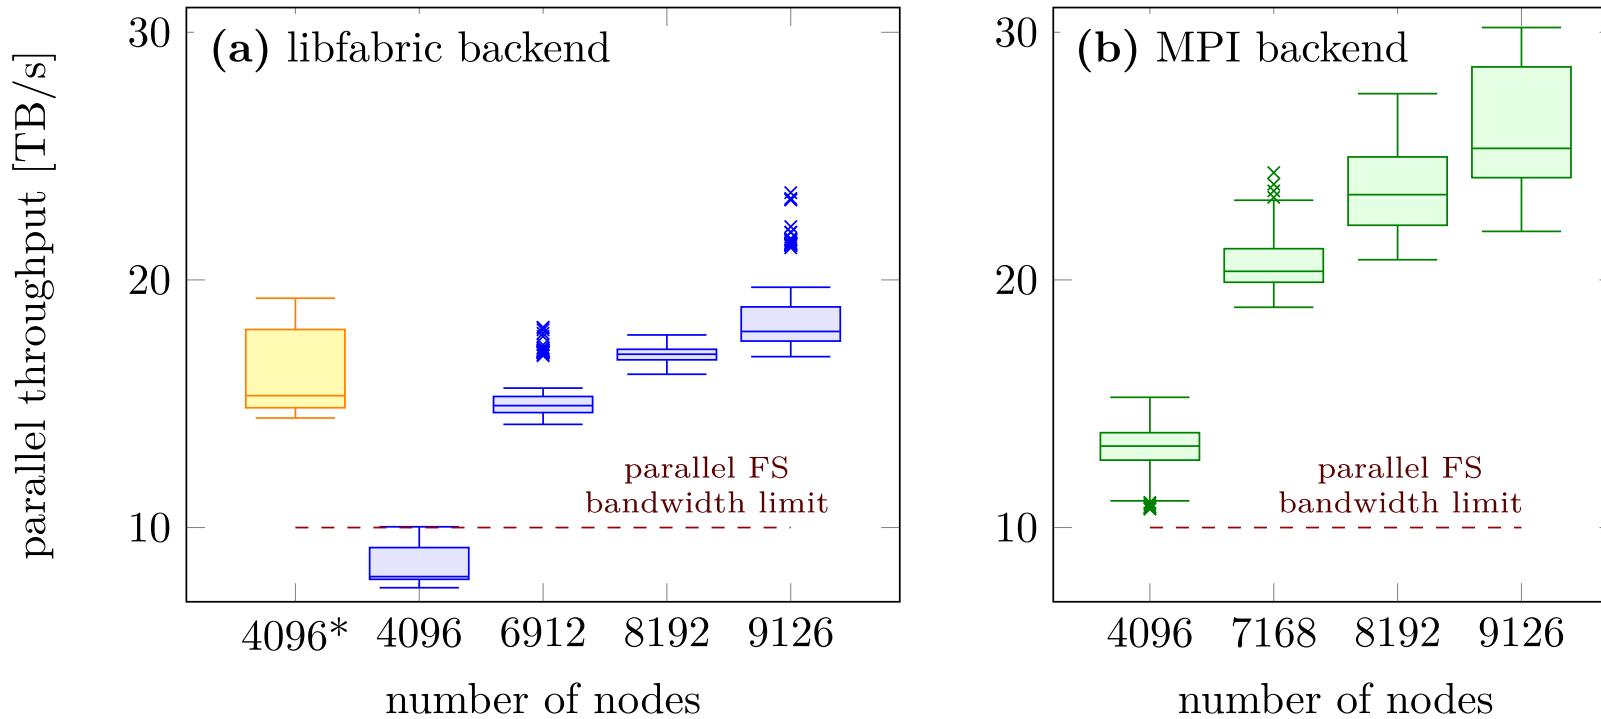
Dashed lines:

TCP/sockets



Setup: Couple PIConGPU with a scattering code (GAPD)
exchange particle data only

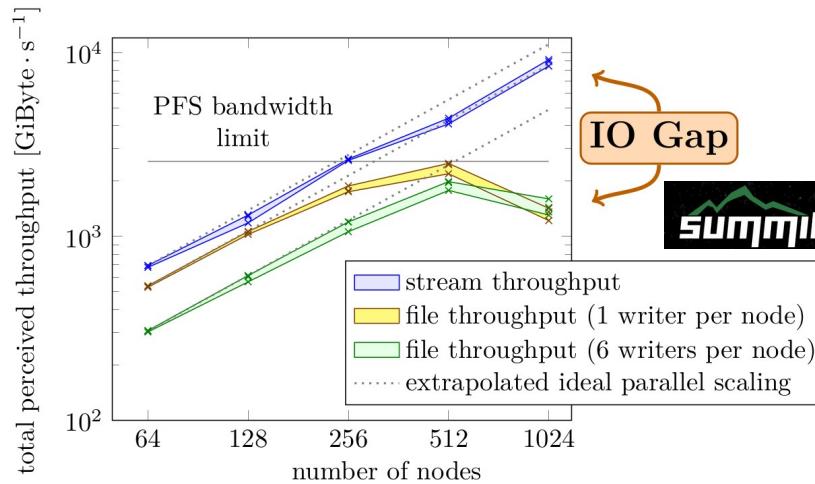
Break through Filesystem Bandwidth with Streaming



Full scale of Frontier:
 Leaving behind us
 the theoretical bandwidth
 of the parallel filesystem



Performance: Data Layouts and no-file I/O



Streaming Data Pipelines:

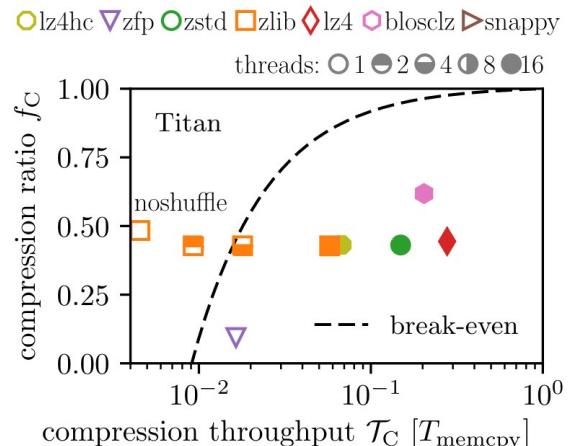
[DOI:10.1007/978-3-030-96498-6_6](https://doi.org/10.1007/978-3-030-96498-6_6)

by F Poeschel, A Huebl et al., SMC21 (2022)

Online Data Layout Reorganization:

[DOI:10.1109/TPDS.2021.3100784](https://doi.org/10.1109/TPDS.2021.3100784)

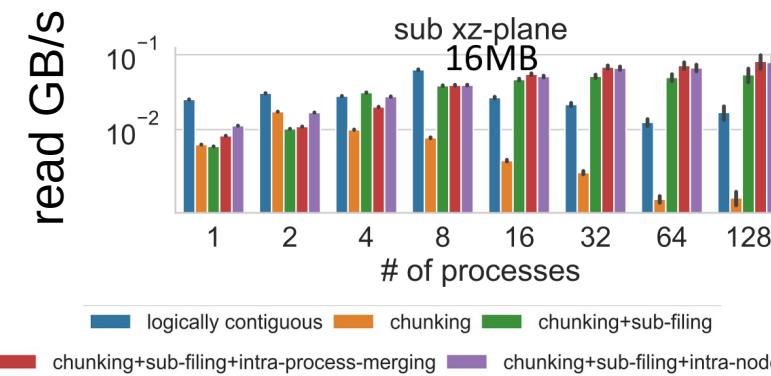
by L Wan, A Huebl et al., TPDS (2021)



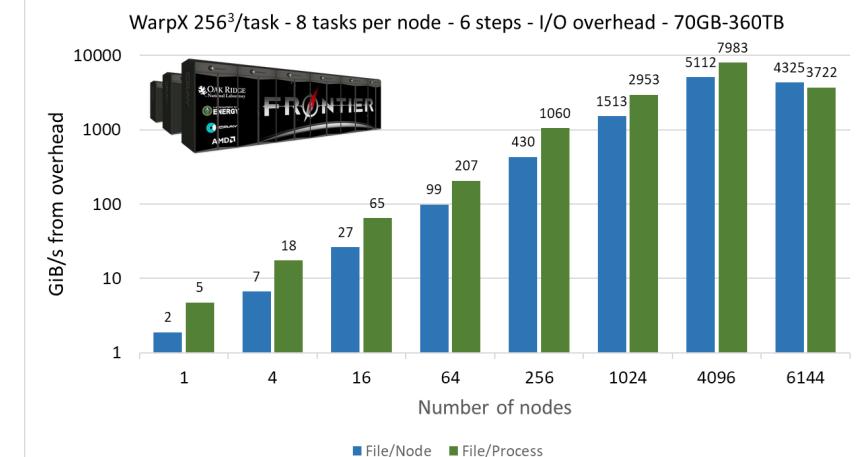
Fast Compressors Needed:

[DOI:10.1007/978-3-319-67630-2_2](https://doi.org/10.1007/978-3-319-67630-2_2)

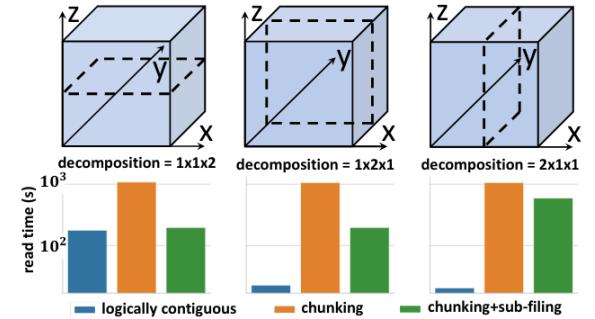
by A Huebl et al., ISC DRBSD-1 (2017)



ADIOS openPMD-api w/ WarpX



>5.5TB/s FS BW: two-tier lustre w/ high-performance storage & progressive files



Impact of decomposition schemes when reading

Hands-On: **openPMD-api:** streaming I/O

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the instructions inside `src/visualize.py` (open as a Jupyter Notebook):



openPMD powered Projects and Users

Documents:

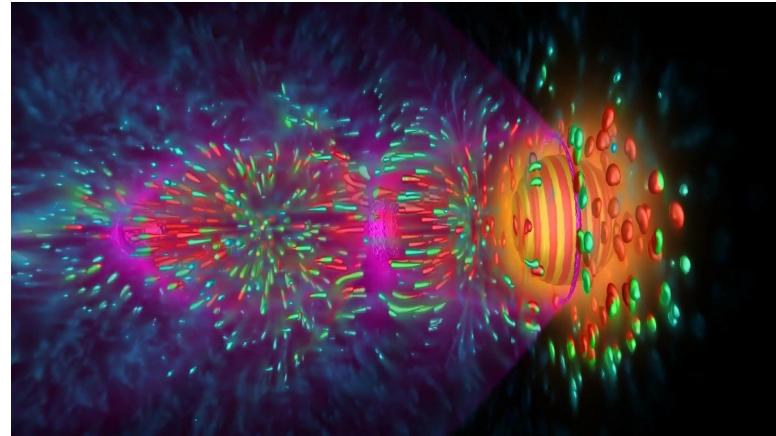
- **openPMD standard** (1.0.0, 1.0.1, 1.1.0)
the underlying file markup and definition
 A Huebl et al., doi: 10.5281/zenodo.33624

Scientific Simulations:

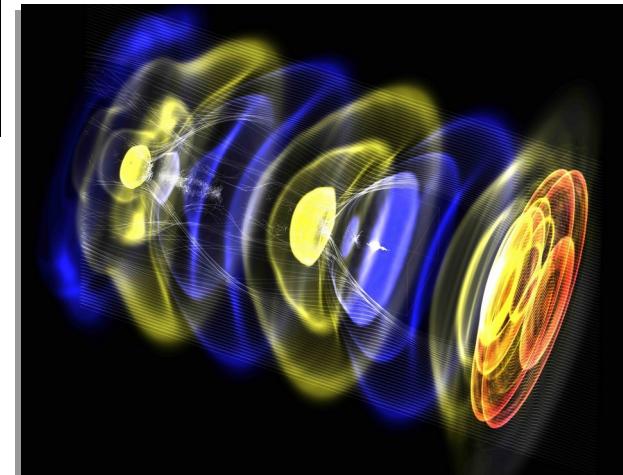
- **PIConGPU** (HZDR)
electro-dynamic particle-in-cell code
 maintainers: R Widera, S Bastrakov, A Debus et al.
- **WarpX** (LBNL, LLNL)
electro-dynamic/static particle-in-cell code
 maintainers: JL Vay, D Grote, R Lehe, A Huebl et al.
- **FBPIC** (LBNL, DESY)
spectral, fourier-bessel particle-in-cell code
 maintainers: R Lehe, M Kirchen et al.
- **SimEx Platform** (EUCALL, European XFEL)
simulation of advanced photon experiments
 maintainer: C Fortmann-Grote

Language Binding:

- **openPMD-api** (HZDR, CASUS, LBNL)
reference API for openPMD data handling
 maintainers: A Huebl, J Gu, F Poeschel et al.



WarpX
 PI: Jean-Luc Vay/LBNL



see also: <https://github.com/openPMD/openPMD-projects>



openPMD powered Projects and Users

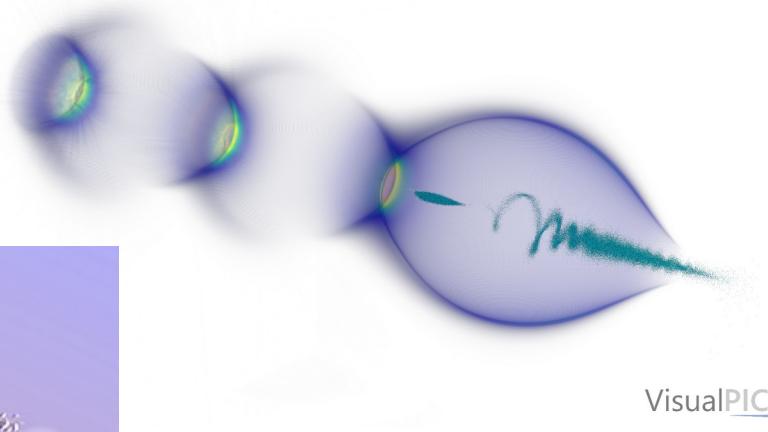
Documents:

- **openPMD standard** (1.0.0, 1.0.1, 1.1.0)
the underlying file markup and definition
 A Huebl et al., doi: 10.5281/zenodo.33624

Language Binding:

- **openPMD-api** (HZDR, CASUS, LBNL)
reference API for openPMD data handling
 maintainers: A Huebl, J Gu, F Poeschel et al.

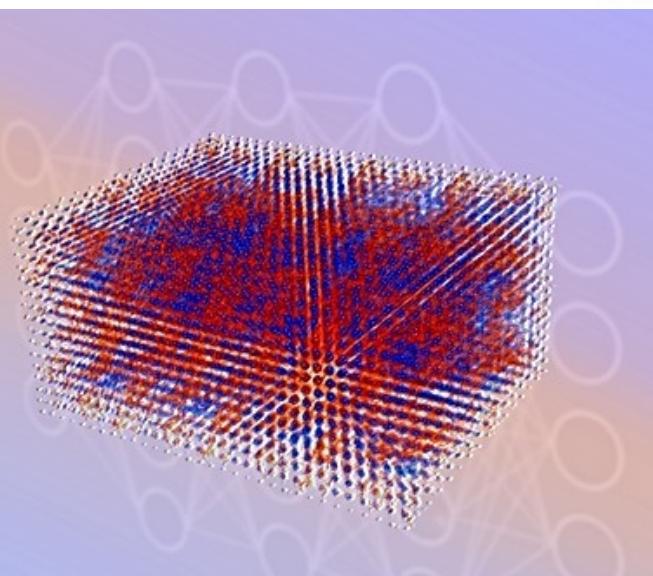
HiPACE++ → VisualPIC
 Credit: M.Thévenet & A. Ferran Pousa (DESY)



VisualPIC

- **Wake-T** (DESY)
fast particle-tracking code for plasma-based accelerators
 maintainer: A Ferran Pousa
- **HiPACE++** (DESY, LBNL)
3D GPU-capable quasi-static PIC code for plasma accel.
 maintainers: M Thevenet, S Diederichs, A Huebl
- **Bmad** (Cornell)
library for charged-particle dynamics simulations
 maintainers: D Sagan et al.
- **MALA** (CASUS, SNL)
ML models that replace DFT calculations in materials science
 maintainers: Attila Cangi & Sivasankaran Rajamanickam
- and more...

see also: <https://github.com/openPMD/openPMD-projects>



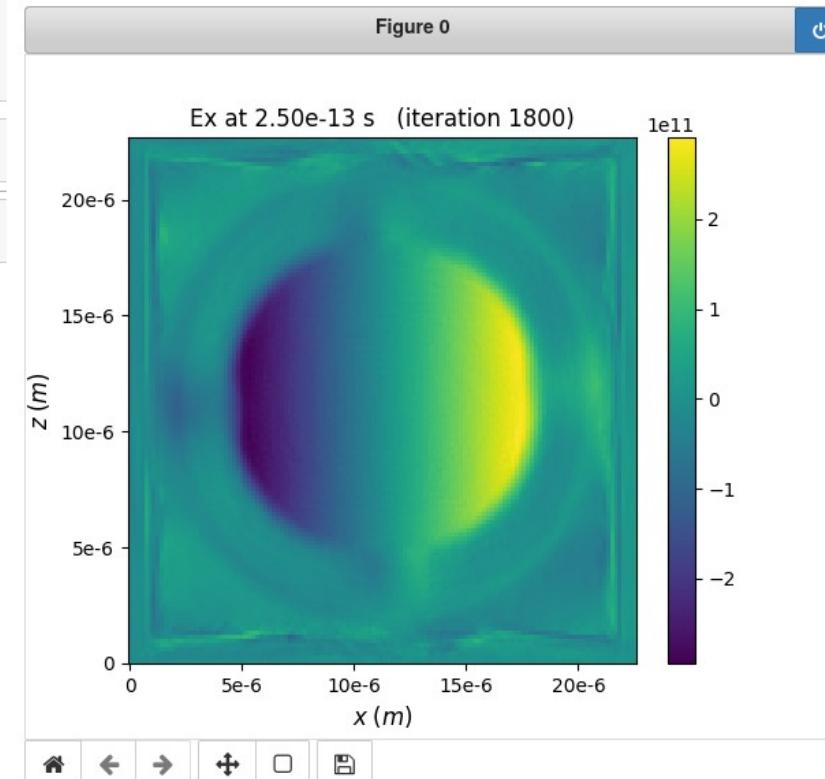
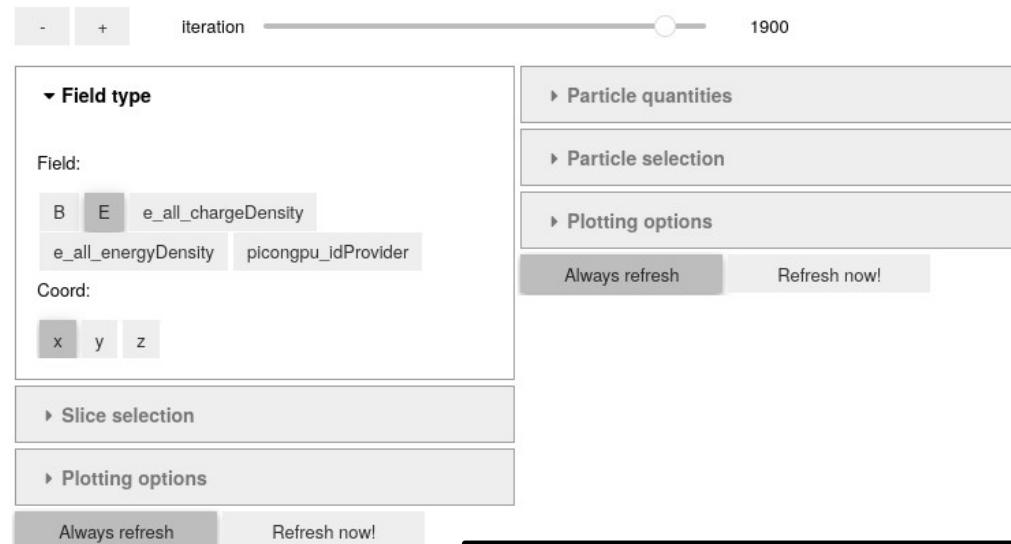
MALA → ParaView
 Credit: A. Cangi (CASUS)

Analysis and Visualization

```
In [1]: import numpy as np
%matplotlib notebook
# or '%matplotlib inline' for non-interactive plots
# or '%matplotlib widget' when using JupyterLab (github.com/matplotlib/jupyter-matplotlib)
import matplotlib.pyplot as plt
from openpmd_viewer import OpenPMDTimeSeries
```

```
In [2]: # Replace the string below, to point to your data
ts = OpenPMDTimeSeries('/home/franzpoeschel/singularity_build/pic_run/openPMD')
```

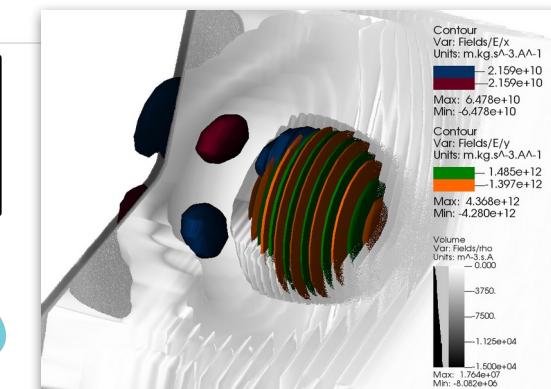
```
In [3]: # Interactive GUI
ts.slider()
```



openPMD/openPMD-viewer



VI-SIT



Standardization of data
 → integration into modern scientific compute workflows

RAPIDS


open
PMD



DASK

Paraview



Hands-On:
openPMD-viewer:
visualizing data written by
a PIConGPU LaserWakefield simulation

Module environment at `/project/project_465001310/workshop_software/env.sh`

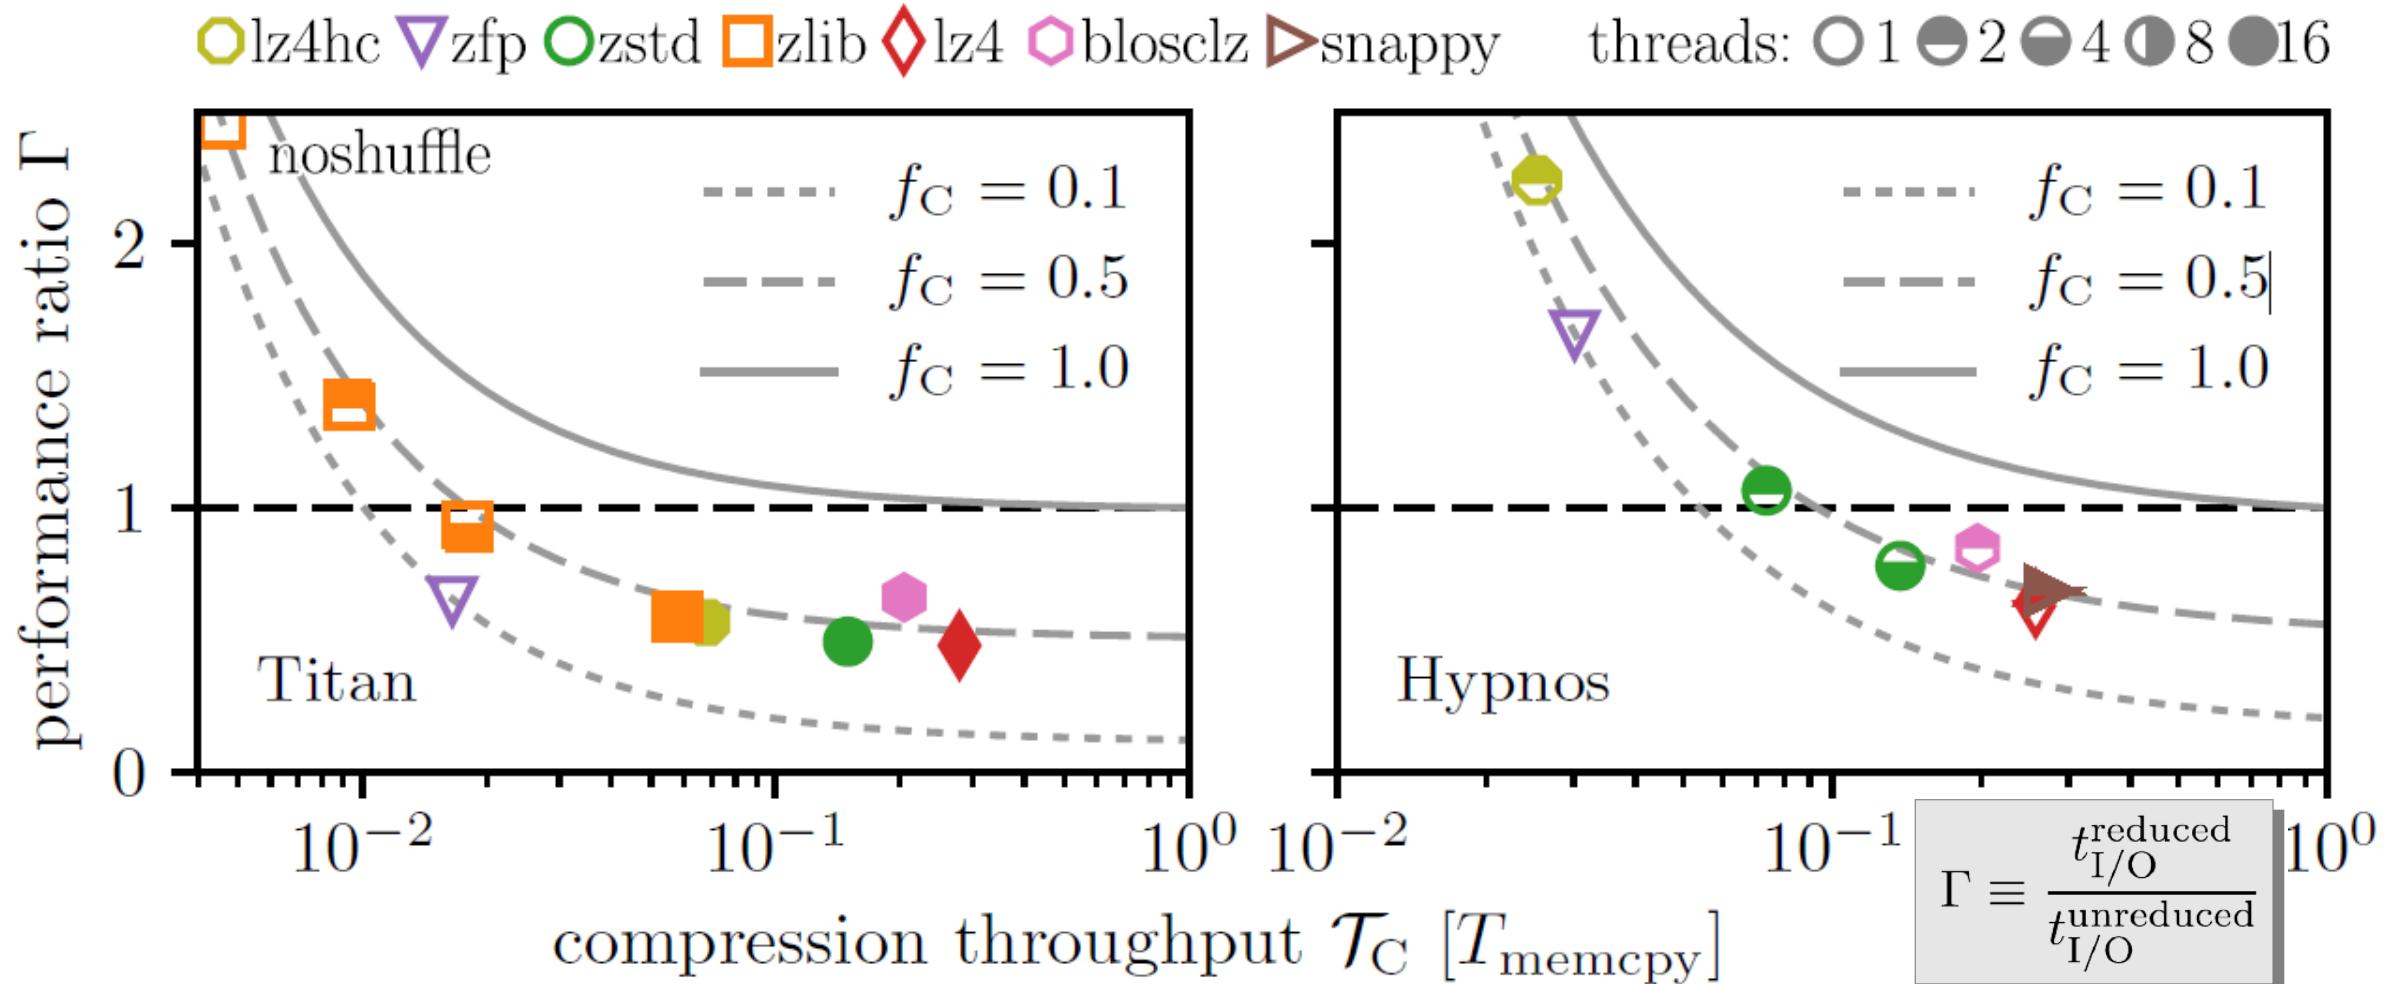
Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the instructions inside `next_steps.md`:



Just compress the data and everything will be fine ...?



Hands-On: **compression**

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

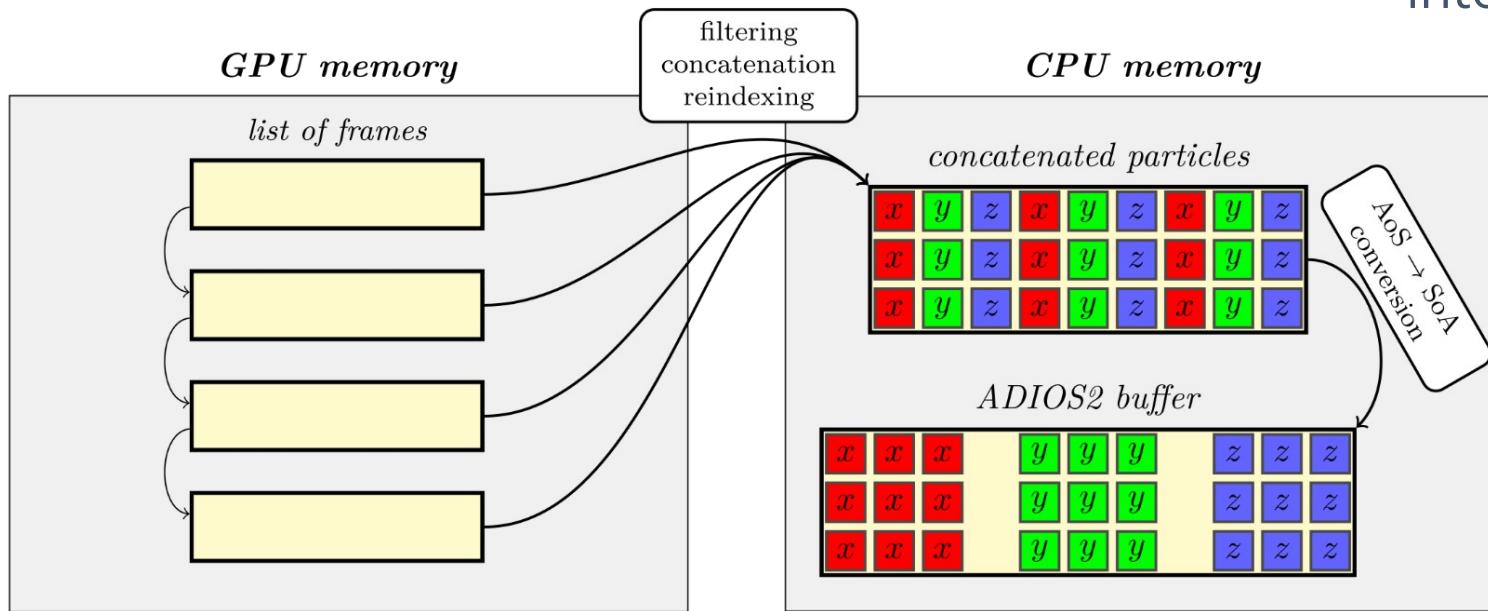
Continue with the instructions in the openPMD-viewer Notebook



Data path in a realistic application

- In real-world applications, data is often kept in a custom, algorithm-driven format
- Reorganization needed before output

- However: ADIOS2 (optionally) buffers data before output for I/O performance
- Few big operations better than many small operations
- Elide one data copy by writing directly into those internal buffers



Example:
 Transforming
 PICoNGPU particle data
 for output

Advanced Hands-On: **openPMD-api:** Span API

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the instructions inside `src/openPMDOutput.hpp`



Modeling particle data

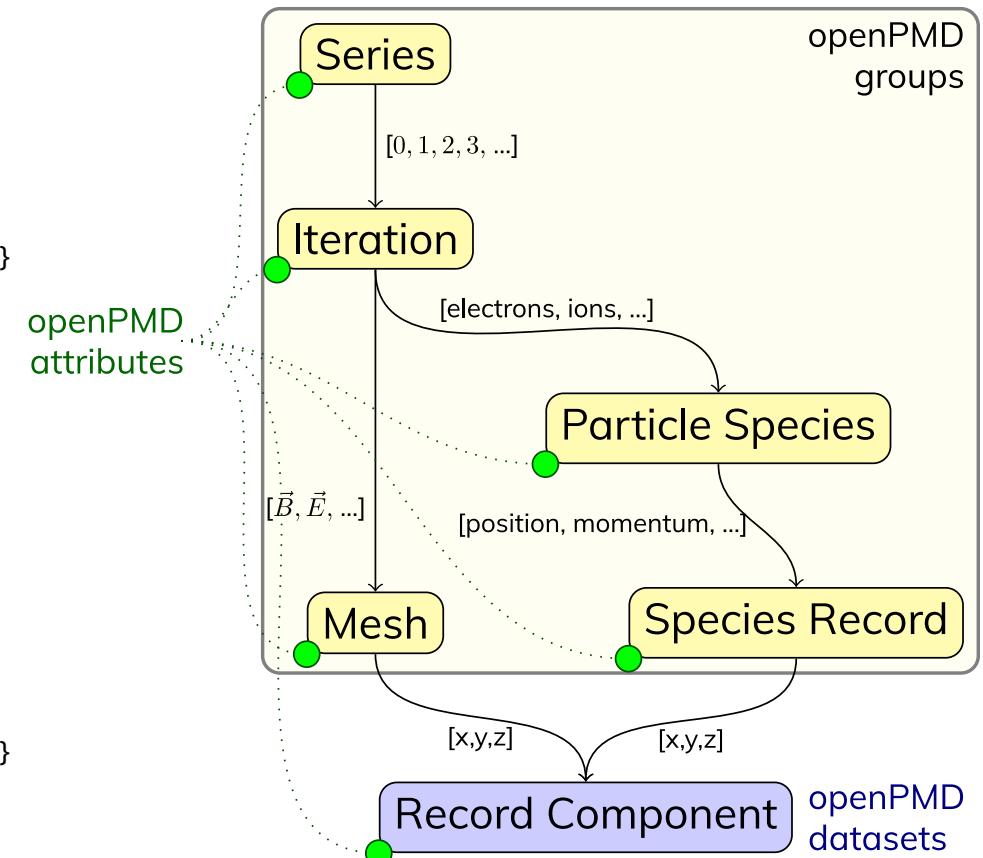
```

uint32_t /data/particles/e/position/macroWeighted      attr   = 0
float    /data/particles/e/position/timeOffset        attr   = 0
double   /data/particles/e/position/unitDimension     attr   = {1, 0, 0, 0, 0, 0, 0}
double   /data/particles/e/position/weightingPower    attr   = 0
float   /data/particles/e/position/x                 {109134176} = 0 / 0
double   /data/particles/e/position/x/unitSI          attr   = 1.772e-07
float   /data/particles/e/position/y                 {109134176} = 0 / 0
double   /data/particles/e/position/y/unitSI          attr   = 4.43e-08
float   /data/particles/e/position/z                 {109134176} = 0 / 0
double   /data/particles/e/position/z/unitSI          attr   = 1.772e-07

uint32_t /data/particles/e/positionOffset/macroWeighted attr   = 0
float    /data/particles/e/positionOffset/timeOffset    attr   = 0
double   /data/particles/e/positionOffset/unitDimension attr   = {1, 0, 0, 0, 0, 0, 0}
double   /data/particles/e/positionOffset/weightingPower attr   = 0
int32_t /data/particles/e/positionOffset/x           {109134176} = 0 / 0
double   /data/particles/e/positionOffset/x/unitSI     attr   = 1.772e-07
int32_t /data/particles/e/positionOffset/y           {109134176} = 0 / 0
double   /data/particles/e/positionOffset/y/unitSI     attr   = 4.43e-08
int32_t /data/particles/e/positionOffset/z           {109134176} = 0 / 0
double   /data/particles/e/positionOffset/z/unitSI     attr   = 1.772e-07

float   /data/particles/e/weighting                  {109134176} = 0 / 0
uint32_t /data/particles/e/weighting/macroWeighted    attr   = 1
float    /data/particles/e/weighting/timeOffset        attr   = 0
double   /data/particles/e/weighting/unitDimension     attr   = {0, 0, 0, 0, 0, 0, 0}
double   /data/particles/e/weighting/unitSI            attr   = 1
double   /data/particles/e/weighting/weightingPower    attr   = 1
  
```

Let's have a detailed look
at the electron species
written by PICoGPU...



Modeling particle data: mandatory records, vector vs. scalar

```

uint32_t /data/particles/e/position/macroWeighted      attr   = 0
float    /data/particles/e/position/timeOffset        attr   = 0
double   /data/particles/e/position/unitDimension     attr   = {1, 0, 0, 0, 0, 0, 0}
double   /data/particles/e/position/weightingPower    attr   = 0
float   /data/particles/e/position/x                 {109134176} = 0 / 0
double   /data/particles/e/position/x/unitSI          attr   = 1.772e-07
float   /data/particles/e/position/y                 {109134176} = 0 / 0
double   /data/particles/e/position/y/unitSI          attr   = 4.43e-08
float   /data/particles/e/position/z                 {109134176} = 0 / 0
double   /data/particles/e/position/z/unitSI          attr   = 1.772e-07

uint32_t /data/particles/e/positionOffset/macroWeighted attr   = 0
float    /data/particles/e/positionOffset/timeOffset    attr   = 0
double   /data/particles/e/positionOffset/unitDimension attr   = {1, 0, 0, 0, 0, 0, 0}
double   /data/particles/e/positionOffset/weightingPower attr   = 0
int32_t /data/particles/e/positionOffset/x           {109134176} = 0 / 0
double   /data/particles/e/positionOffset/x/unitSI     attr   = 1.772e-07
int32_t /data/particles/e/positionOffset/y           {109134176} = 0 / 0
double   /data/particles/e/positionOffset/y/unitSI     attr   = 4.43e-08
int32_t /data/particles/e/positionOffset/z           {109134176} = 0 / 0
double   /data/particles/e/positionOffset/z/unitSI     attr   = 1.772e-07

float   /data/particles/e/weighting                  {109134176} = 0 / 0
uint32_t /data/particles/e/weighting/macroWeighted    attr   = 1
float    /data/particles/e/weighting/timeOffset        attr   = 0
double   /data/particles/e/weighting/unitDimension     attr   = {0, 0, 0, 0, 0, 0, 0}
double   /data/particles/e/weighting/unitSI            attr   = 1
double   /data/particles/e/weighting/weightingPower    attr   = 1

```

position: mandatory record for particles

position is a **vector quantity**, i.e. has x/y/z struct-of-array-like layout:

3 vectors of same length

positionOffset: also mandatory specifies the base for the position used for numerical reasons

weighting: custom record in PICoNGPU to model macro particles scalar quantity, x/y/z is skipped

Modeling particle data: constant components

```

string /data/particles/e/currentDeposition           attr = "Esirkepov"
string /data/particles/e/particleInterpolation       attr = "uniform"
string /data/particles/e/particlePush                attr = "Boris"
double /data/particles/e/particleShape               attr = 2
string /data/particles/e/particleSmoothing          attr = "none"

int32_t /data/particles/e/charge/macroWeighted      attr = 0
uint64_t /data/particles/e/charge/shape             attr = {109134176}
double /data/particles/e/charge/timeOffset           attr = 0
double /data/particles/e/charge/unitDimension        attr = {0, 0, 1, 1, 0, 0, 0}
double /data/particles/e/charge/unitSI               attr = 1.11432e-15
double /data/particles/e/charge/value              attr = -0.00014378
double /data/particles/e/charge/weightingPower       attr = 1

int32_t /data/particles/e/mass/macroWeighted        attr = 0
uint64_t /data/particles/e/mass/shape              attr = {109134176}
double /data/particles/e/mass/timeOffset             attr = 0
double /data/particles/e/mass/unitDimension          attr = {0, 1, 0, 0, 0, 0, 0}
double /data/particles/e/mass/unitSI                 attr = 6.33564e-27
double /data/particles/e/mass/value                attr = 0.00014378
double /data/particles/e/mass/weightingPower          attr = 1

uint32_t /data/particles/e/momentum/macroWeighted attr = 1
float /data/particles/e/momentum/timeOffset          attr = 0
double /data/particles/e/momentum/unitDimension       attr = {1, 1, -1, 0, 0, 0, 0}
double /data/particles/e/momentum/weightingPower      attr = 1
float /data/particles/e/momentum/x                  attr = {109134176} = 0 / 0
double /data/particles/e/momentum/x/unitSI            attr = 1.89938e-18
float /data/particles/e/momentum/y                  attr = {109134176} = 0 / 0
double /data/particles/e/momentum/y/unitSI            attr = 1.89938e-18
float /data/particles/e/momentum/z                  attr = {109134176} = 0 / 0
double /data/particles/e/momentum/z/unitSI            attr = 1.89938e-18
  
```

particle attributes defined by the ED-PIC extension

charge and mass:

are fixed for one particle species,
i.e. record has one constant value

→ define via attributes **shape** and **value**
instead of writing the whole array
scalar quantities as well

momentum:

again a non-constant vector quantity

Modeling particle data: particle patches

```

uint64_t /data/1000/particles/e/particlePatches/offset/x      {8}
(0)   0 0 0 0 96 96 96 96

uint64_t /data/1000/particles/e/particlePatches/offset/y      {8}
(0)   0 512 1024 1536 0 512 1024 1536

uint64_t /data/1000/particles/e/particlePatches/offset/z      {8}
(0)   0 0 0 0 0 0 0 0

uint64_t /data/1000/particles/e/particlePatches/extent/x      {8}
(0)   96 96 96 96 96 96 96 96

uint64_t /data/1000/particles/e/particlePatches/extent/y      {8}
(0)   512 512 512 0 512 512 512 0

uint64_t /data/1000/particles/e/particlePatches/extent/z      {8}
(0)   192 192 192 192 192 192 192 192

uint64_t /data/1000/particles/e/particlePatches/numParticlesOffset {8}
(0)   0 16916087 35692896 54567264 54567264 71482787 90259808 109134176

uint64_t /data/1000/particles/e/particlePatches/numParticles      {8}
(0)   16916087 18776809 18874368 0 16915523 18777021 18874368 0

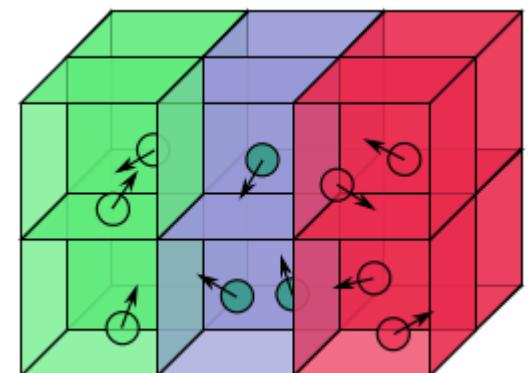
```

green: GPU 0

- Common case:** Each GPU computes one sub-area of the global simulation space
- Particles are written in patches, each GPU's patch is spatially bounded
- offset/extent define the box wherein the particles are located
 - numParticlesOffset/numParticles define the patch's position in the list

e.g.: GPU 0 computes the box (0, 0, 0)–(96, 512, 192) (in simulation units), its electron data is found from index 0 to 16916087

Note: extent is always relative to the offset



Advanced Hands-On: **openPMD-api:** Python API: constant components

Module environment at `/project/project_465001310/workshop_software/env.sh`

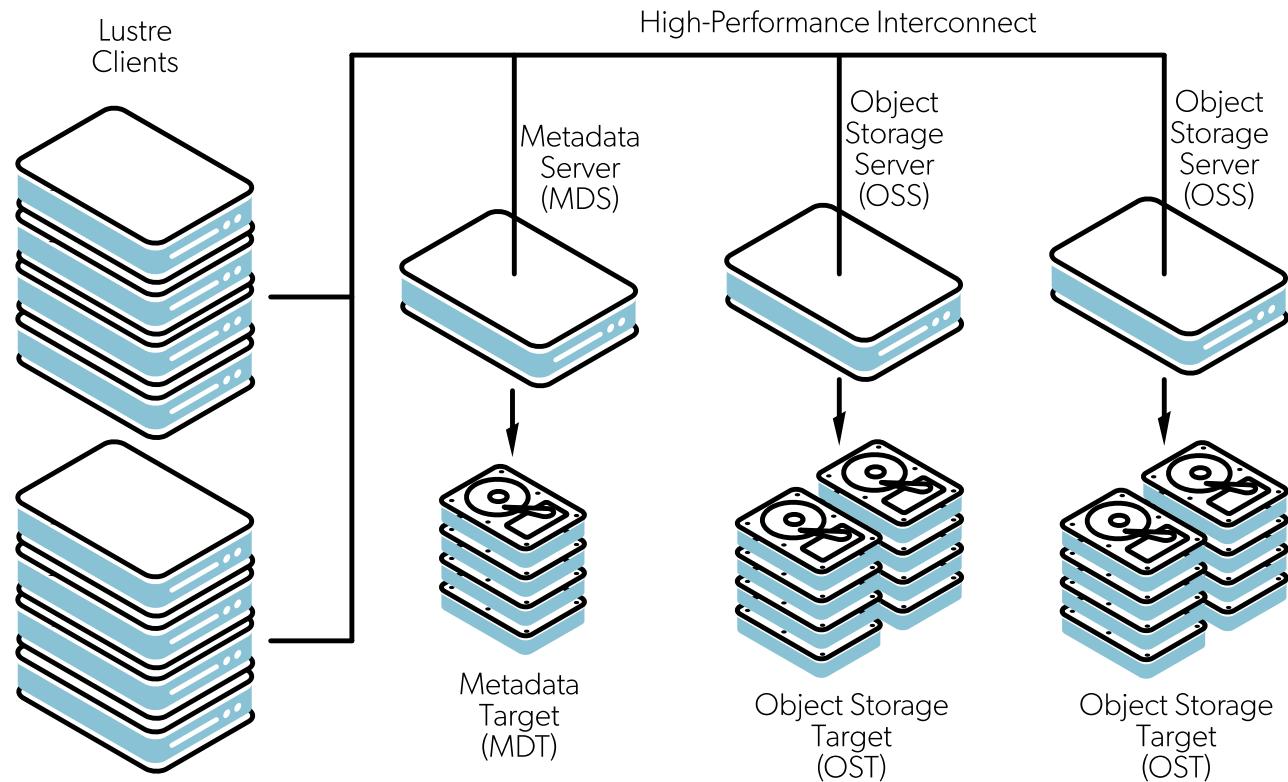
Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the TODO comments inside `write_parallel.py`



Using parallel filesystems



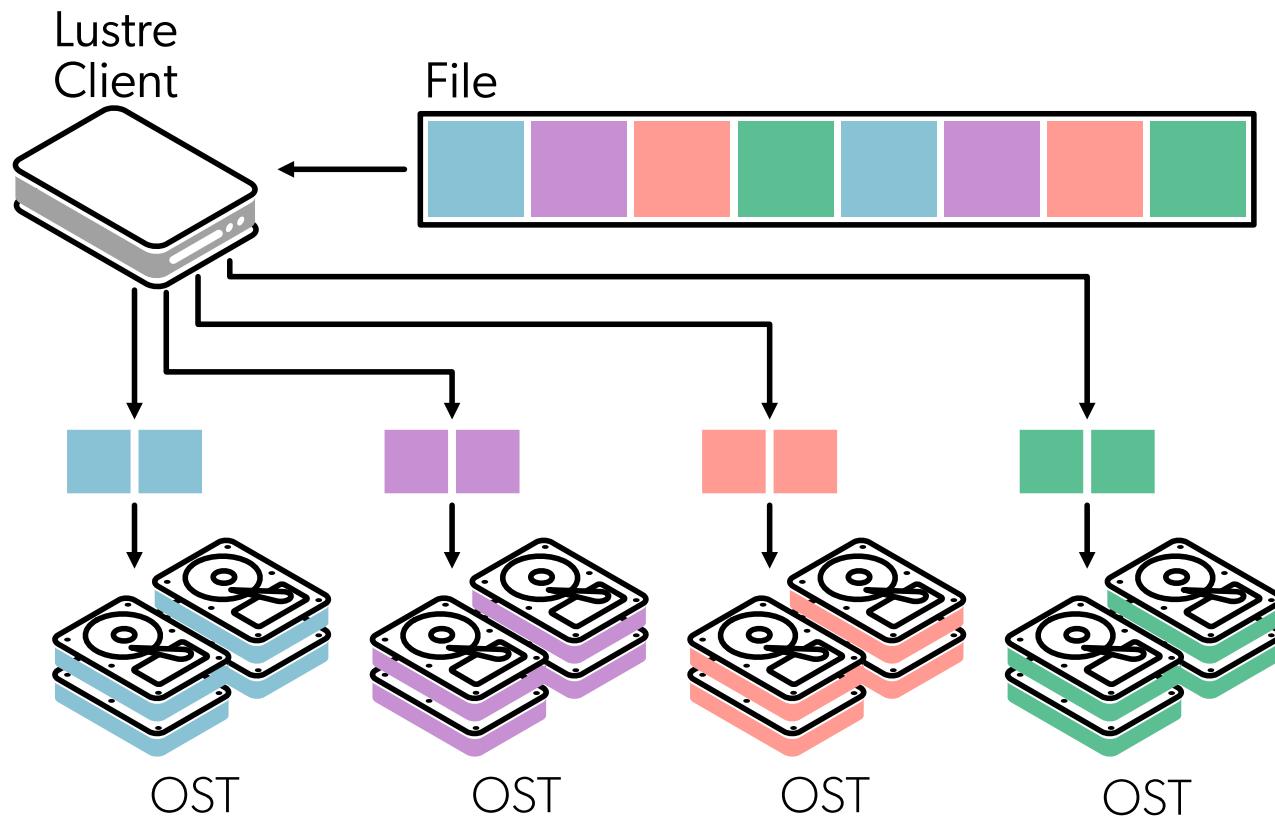
Lumi uses Lustre,
a popular parallel filesystem

- Parallel filesystems are built from distributed hardware
- One or more metadata servers manage one or more metadata targets
- Many object storage servers manage each one or more object storage targets
- Key to I/O performance:
Use the parallel resources
(but mind that they are shared with other users)

Image from

<https://docs.lumi-supercomputer.eu/storage/parallel-filesystems/lustre/>

File striping - an approach at parallel I/O performance



Write file in parallel by striping it across multiple OSTs

- **stripe_size** sets the homogeneous length of each stripe
- **stripe_count** sets the number of used OSTs

Main use:

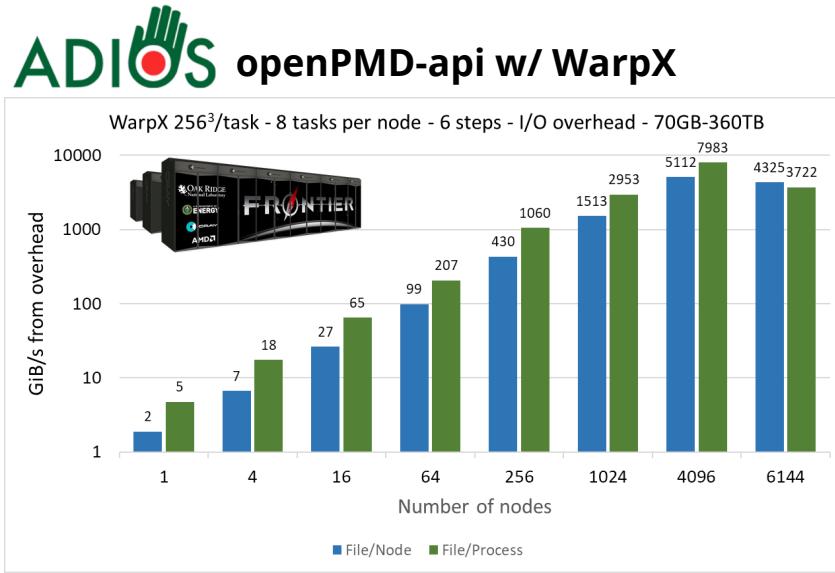
- Improving speed of **serial I/O**
- Improving speed of parallel I/O to a **single file**
- At low scale: Improving speed of parallel I/O to **multiple files**

Better: Write **multiple files** from the start already → **Subfiling**

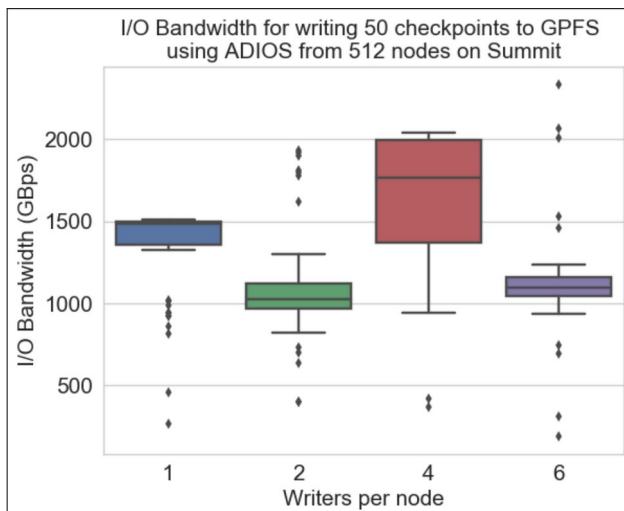
Image from

<https://docs.lumi-supercomputer.eu/storage/parallel-filesystems/lustre/>

Subfiling - key for performance at full scale



>5.5TB/s FS BW: two-tier lustre w/ high-performance storage & progressive files



L. Wan et al.: Data Management Challenges of Exascale Scientific Simulations:
A Case Study with the Gyrokinetic Toroidal Code and ADIOS

Subfiling is NOT independent I/O

- Mapping processes to files N:1 scales badly
- Mapping processes to files N:N scales up to a mediocre size, then overloads the filesystem
- Subfiling: N:M mapping where M≤N
- Many different aggregation schemes in ADIOS2 and HDF5, one size does not fit all

ADIOS2: Aggregate to one writer per node by default

HDF5: Complex N→N aggregation scheme due to structured output to disk

Advanced Hands-On: **openPMD-api:** Python API: Parallel Writing

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the TODO comments inside `write_parallel.py`



Where to find us

→ head to <https://github.com/openPMD/>



openPMD

Open Standard for Particle-Mesh Data Files

🔗 <https://www.openPMD.org> 📩 axelhuebl@lbl.gov

Repositories 17 Packages People 50 Teams 5 Projects

Pinned repositories

[openPMD-standard](#)
Open Standard for Particle-Mesh Data Files
41 stars, 17 forks

[openPMD-projects](#)
Overview on Projects around openPMD
4 stars, 4 forks

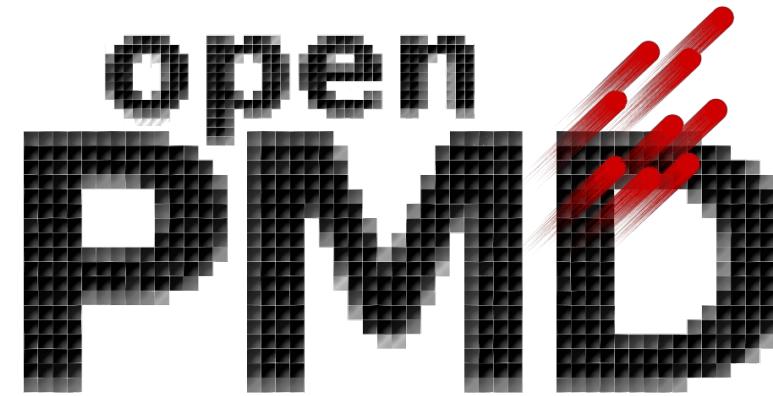
[openPMD-viewer](#)
Python visualization tools for openPMD files
35 stars, 26 forks

[openPMD-api](#)
C++ & Python API for Scientific I/O
C++ language, 55 stars, 30 forks

[openPMD-visit-plugin](#)
Plugin allowing VisIt to read openPMD files
C language, 8 stars, 3 forks

[openPMD-example-datasets](#)
HDF5 Example Files
Python language, 5 stars, 1 fork

...and of course <https://openpmd-api.readthedocs.io/>



The **openPMD standard** is co-authored by [Axel Huebl](#), [Rémi Lehe](#), Jean-Luc Vay, David P. Grote, Ivo F. Sbalzarini, Stephan Kuschel, David Sagan, Frédéric Pérez, Fabian Koller, [Franz Poeschel](#), Carsten Fortmann-Grote, Ángel Ferran Pousa, Juncheng E, [Maxence Thévenet](#), and Michael Bussmann.

The authors are thankful for the **community contributions** to libraries, software ecosystem, user support, review and integrations. Particularly, thank you to Yaser Afshar, Lígia Diana Amorim, James Amundson, Weiming An, Igor Andriyash, Ksenia Bastrakova, [Jean Luca Bez](#), Richard Briggs, Heiko Burau, Jong Choi, Ray Donnelly, Dmitry Ganyushin, Marco Garten, Lixin Ge, Berk Geveci, Daniel Grassinger, Alexander Grund, [Junmin Gu](#), Marc W. Guetg, Ulrik Günther, Sören Jalas, Manuel Kirchen, John Kirkham, Scott Klasky, Noah Klemm, Fabian Koller, Mathieu Lobet, Christopher Mayes, Ritiek Malhotra, Paweł Ordyna, Richard Pausch, [Norbert Podhorszki](#), David Pugmire, Felix Schmitt, [Erik Schnetter](#), Dominik Stańczak, Klaus Steiniger, Michael Sippel, Frank Tsung, Lipeng Wan, René Widera, and Erik Zenker!