

Heat Equation Simulation using Alpaka

Alpaka Team

HZDR

September 5, 2024

Overview

- 1 Introduction
- 2 Domain and Stencil
- 3 Parallelization
- 4 Simulation Setup using Alpaka
- 5 WorkDiv: How to parallelize?
- 6 Execution
- 7 Optimization
- 8 Conclusion

The Heat Equation

- The heat equation models the Heat Diffusion over time in a given medium.

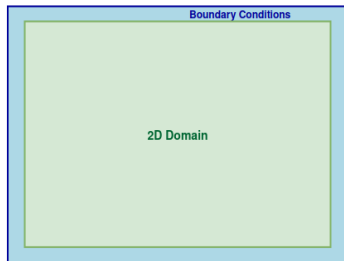
$$\frac{\partial u(x, y, t)}{\partial t} = \alpha \left(\frac{\partial^2 u(x, y, t)}{\partial x^2} + \frac{\partial^2 u(x, y, t)}{\partial y^2} \right)$$

Difference approximations for Time and Spatial Derivatives:

$$\left. \frac{\partial u(x, y, t)}{\partial t} \right|_{t=t^n} \approx \frac{u_{i,j}^{n+1} - u_{i,j}^n}{\Delta t} \quad \left. \frac{\partial^2 u(x, y, t)}{\partial x^2} \right|_{x=x_i, y=y_j} \approx \frac{u_{i+1,j}^n - 2u_{i,j}^n + u_{i-1,j}^n}{\Delta x^2}$$


- Resulting difference equation:

$$u_{i,j}^{n+1} = u_{i,j}^n + \alpha \Delta t \left(\frac{u_{i+1,j}^n - 2u_{i,j}^n + u_{i-1,j}^n}{\Delta x^2} + \frac{u_{i,j+1}^n - 2u_{i,j}^n + u_{i,j-1}^n}{\Delta y^2} \right)$$



Parallel Heat Equation Solution

- **Data Parallelism:** Each point on the grid can be updated independently based on its neighbors, enabling parallel computation.
- **Stencil Operations:** Stencil is a core computational pattern in PDE solvers. Updates a grid point in time using its immediate neighbors (left, right, up, down) according to the difference equation. A 5-point stencil is needed.



- **Halo Region for BC:** A layer of grid cells surrounding the problem domain for Boundary Conditions.
 - Facilitates stencil operations at the boundaries of subdomains.

Calculation of $u_{i,j}^{n+1}$ from $u_{i,j}^n$

- Each kernel execution by alpaka calculates $u_{i,j}^{n+1}$ using $u_{i,j}^n$
- Each heat point is separately calculated by a thread using frobenious inner product
- Calculation of $u_{0,0}^{n+1}$ by a single thread

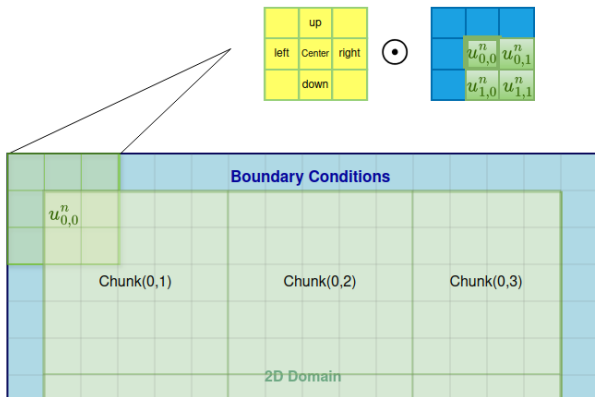


Figure: First thread calculates $u_{0,0}^{n+1}$ using frobenious inner product of 3x3 matrices

- Second thread calculates $u_{0,1}^{n+1}$

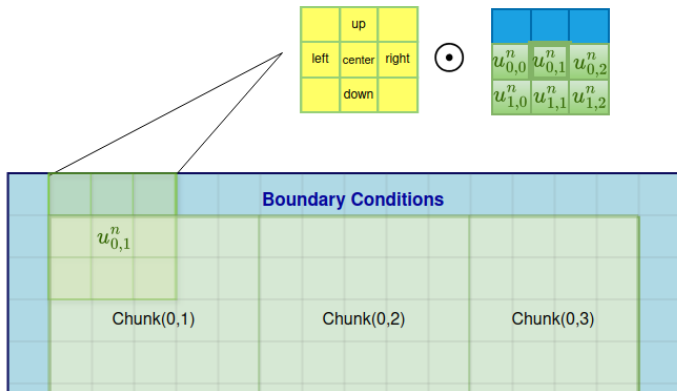




Figure: Second thread calculates $u_{0,1}^{n+1}$ using

Chunks in Parallel Grid Computations

- **Chunk:** Subdomains needed for block level parallelisation
- **Halo Region around chunk:** A layer of grid cells surrounding the subdomains. In order to use the heat value beside the current chunk
- **Halo Size:** Typically 1 for a 5-point stencil.



Chunk Definition (in detail?)

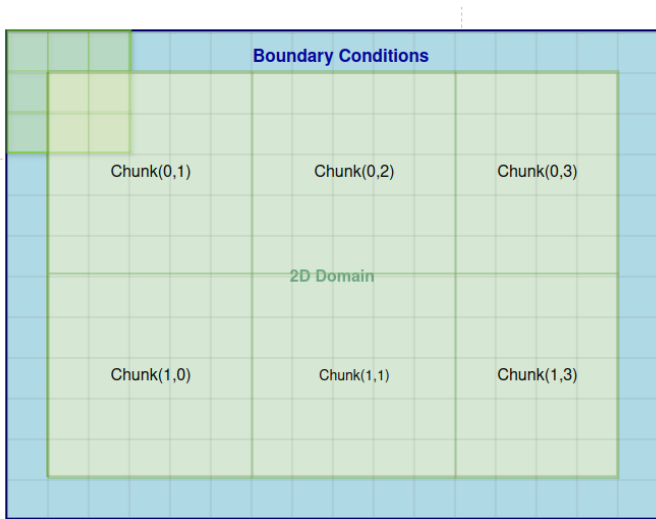


Figure: Heat values corresponding to each chunk is updated by a block

Accelerator, Device and Host

Define number of dim and index type

```
1 using Dim = alpaka::DimInt<2u>; // Number of dim: 2 as a type
2 using Idx = std::size_t; // Index type of the threads and buffers
```

Define the accelerator

```
1 // AccGpuCudaRt, AccGpuHipRt, AccCpuThreads, AccCpuSerial,
2 // AccCpuOmp2Threads, AccCpuOmp2Blocks, AccCpuTbbBlocks
3 using Acc = alpaka::AccGpuCudaRt<Dim, Idx>;
4 using DevAcc = alpaka::Dev<Acc>;
```

Select a device from platform of Acc

```
1 auto const platform = alpaka::Platform<Acc>{};
2 auto const devAcc = alpaka::getDevByIdx(platform, 0);
```

Select a host and hosttype to allocate memory for data

```
1 // Get the host device for allocating memory on the host.
2 auto const platformHost = alpaka::PlatformCpu{};
3 auto const devHost = alpaka::getDevByIdx(platformHost, 0);
4 // Host device type is needed, still not known
5 using DevHost = alpaka::DevCpu;
```

Allocate memory at Host and at Device

```

1
2 // Allocate host memory buffers
3 using BufHost = alpaka::Buf<DevHost, DataType, Dim, Idx>;
4 BufHost bufHostA(alpaka::allocBuf<DataType, Idx>(devHost, extent));
5
6 // Fill the host buffers
7 for (Idx i(0); i < numElements; ++i) {
8     bufHostA[i] = randomA;
9 }
10
11 // Allocate buffer on the accelerator
12 using BufAcc = alpaka::Buf<DevAcc, DataType, Dim, Idx>;
13 BufAcc bufAccA(alpaka::allocBuf<DataType, Idx>(devAcc, extent));

```

Need a queue to copy the data to the device

```

1  using Acc = alpaka::AccCpuSerial<Dim, Idx>;
2  auto const platformAcc = alpaka::Platform<Acc>{};
3  auto const devAcc = alpaka::getDevByIdx(platformAcc, 0);
4  // A queue is needed for all acc related operations
5  alpaka::Queue<Acc, alpaka::Blocking> queue(devAcc);
6
7  // Define the 2D extents (dimensions)
8  Vec const extentA(static_cast<Idx>(M), static_cast<Idx>(K));
9
10 // Allocate host memory, the memory size is determined by extent
11 auto bufHostA = alpaka::allocBuf<DataType, Idx>(devHost, extentA
12     );
13
14 // Allocate device memory
15 auto bufDevA = alpaka::allocBuf<DataType, Idx>(devAcc, extentA);
16
17 // Copy data to device, use host buffer and device buffer
18 // Queue must be an accelerator queue
19 alpaka::memcpy(queue, bufDevA, bufHostA);

```


Passing multi dimensional data to the kernel

Multi-dimensional memory allocated in memory uses aligned rows.

Hence, if a pointer of a 2D buffer is passed to the kernel as a pointer; 2 additional values **pitch** and item **data-size** should also be passed.

- Pass 3 variables to kernel: pointer, pitch, and datasize

```
1  template<typename TAcc, typename TDim, typename
2  TIdx>
3  ALPAKA_FN_ACC auto operator()(
4      TAcc const& acc,
5      double const* const uCurrBuf,
6      double* const uNextBuf,
7      alpaka::Vec<TDim, TIdx> const pitchCurr,
8      alpaka::Vec<TDim, TIdx> const pitchNext,
9      ...) const -> void
```



2D Buffer with size 15x7 in memory

- Simple Alternative: Pass an alpaka::mdspan object


```
1  template<typename TAcc, typename TDim, typename TIdx, typename TMdSpan>
2  ALPAKA_FN_ACC auto operator()(
3      TAcc const& acc,
4      TMdSpan uCurrBuf,
5      TMdSpan uNextBuf
6      ...) const -> void
```

WorkDiv-I: Setting workdiv fields directly

```

1  auto blocksPerGrid = alpaka::Vec<Dim, Idx>{M/8, N/128};
2  auto threadsPerBlock = alpaka::Vec<Dim, Idx>{8, 128};
3  auto elementsPerThread = alpaka::Vec<Dim, Idx>{1u, 1u};
4  using WorkDiv = alpaka::WorkDivMembers<Dim, Idx>;
5  auto workDiv = WorkDiv{blocksPerGrid, threadsPerBlock,
    elementsPerThread};

```



WorkDiv-II: Let Alpaka calculate for you

```

1  using Acc = alpaka::AccCpuSerial<Dim, Idx>;
2  auto const devAcc = alpaka::getDevByIdx(platformAcc, 0);
3
4  // Define the 2D extents as {Width, Height} of matrix
5  alpaka::Vec<Dim, Idx> const extentAsThreadsPerGrid(M, N);
6  auto elementsPerThread = alpaka::Vec<Dim, Idx>{1u, 1u};
7  MatrixMulKernel kernel;
8
9  // Let alpaka calculate good block and grid sizes given our full
   problem extent
10 alpaka::KernelCfg<Acc> const kernelCfg = {extentAsThreadsPerGrid,
   elementsPerThread};
11 auto const workDiv = alpaka::getValidWorkDiv<Acc>(kernelCfg, devAcc,
   kernel, // kernel params here
12 );

```

Main Simulation Loop: Leveraging Parallelism

■ Initialization:

- Define the "host device" and "accelerator device". The "Host" and "Device" in short.
- Set initial conditions and boundary conditions.
- Allocate data buffers to host and device.
- Copy data from host to device buffer to pass to the kernel.
- Define parallelisation strategy (chunk size, block size, etc.).

■ Simulation Loop:

- **Step 1:** Execute `StencilKernel` to compute next values.
- **Step 2:** Apply boundary conditions using `BoundaryKernel`.
- **Step 3:** Swap buffers for the next iteration so that calculated $u_{i,j}^{n+1}$ becomes the $u_{i,j}^n$ for the next step.

■ Parallel Efficiency:

- Subdomains are processed in parallel, with halos ensuring data consistency and correct boundary conditions.
- Optimization: Shared memory optimizes memory access within each block.

■ Validation

Executing the Kernel

■ Execution Flow:

- Each kernel is executed on the selected accelerator (e.g., CPU, GPU).
- Halo regions and shared memory are leveraged for optimal parallel performance.

■ Kernel Execution:

```
alpaka::exec<Acc>(
    queue1,
    workDiv_manual,
    stencilKernel,
    uCurrBufAcc.data(),
    uNextBufAcc.data(),
    chunkSize,
    dx, dy, dt);
```

- **Run Example:** Execute for all enabled accelerators (e.g., CUDA, HIP, OpenMP).

Applying Boundary Conditions in Parallel

- **Boundary Kernel:** Ensures correct values at the boundaries of the grid.
- **Challenges in Parallel Computing:**
 - Boundary points might need special handling, particularly when subdomains are processed independently.
 - Halo regions play a crucial role here.

- **Code Snippet:**

```
if (gridBlockIdx[0] == 0) {  
    applyBoundary(globalIdx, chunkSize[1], true);  
}
```

Efficient Stencil Computation with Shared Memory

■ Shared Memory:

- A fast, limited-size memory accessible by all threads within a block.
- Used to store grid points locally, reducing the need to access slower global memory.

■ Benefits:

- Reduces memory latency by storing the working set of data (halo + core) in shared memory.
- Enables efficient data reuse across threads in the same block.

■ Example:

```
auto& sdata =
alpaka::declareSharedVar<double[T_SharedMemSize1D],
__COUNTER__>(acc);
```

■ Synchronization ...I/O:

- Threads in a block must synchronize to ensure all data is loaded into shared memory before computation begins.

Using multiple queues

Results and Performance in Parallel Execution

■ Validation:

- Accuracy of results compared to the analytical solution.
- Performance considerations: Speedup achieved by parallelizing the computation.

■ Output:

- Print whether the results are correct.
- Report on the maximum error.
- Discuss any performance metrics (e.g., execution time).

■ Visual Output (Optional):

- Periodic snapshots of the temperature distribution.

Conclusion: Parallel Techniques for Efficient Simulation

■ Key Takeaways:

- Efficient use of shared memory significantly boosts performance in parallel computations.
- Halo regions are crucial for managing data dependencies in stencil operations.
- The combination of Alpaka's abstraction and careful memory management enables scalable and portable parallel solutions.

Questions?